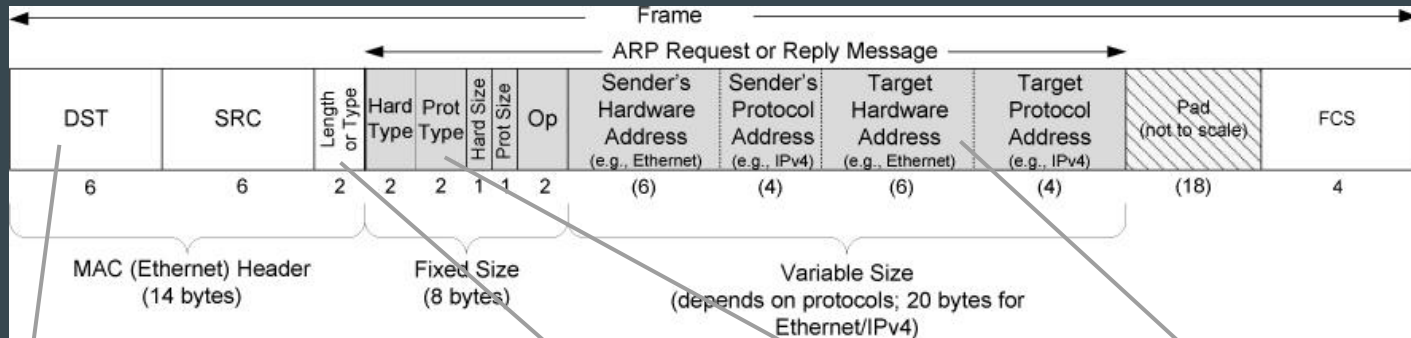


ARP Frame encapsulated in Ethernet Frame



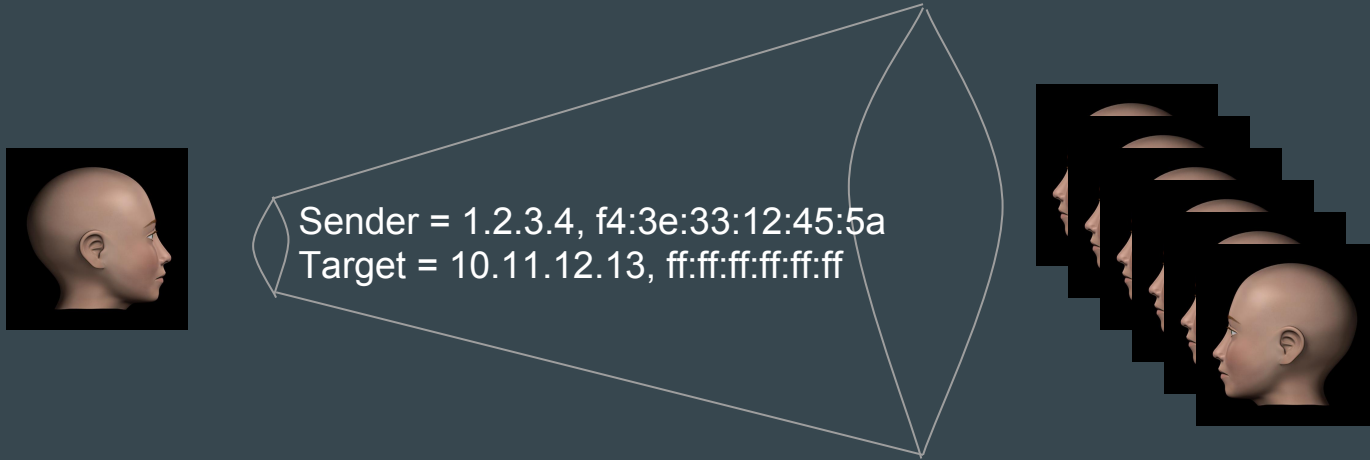
DST = link-layer broadcast address in a request. Unicast in a response.

ARP = 0x0806

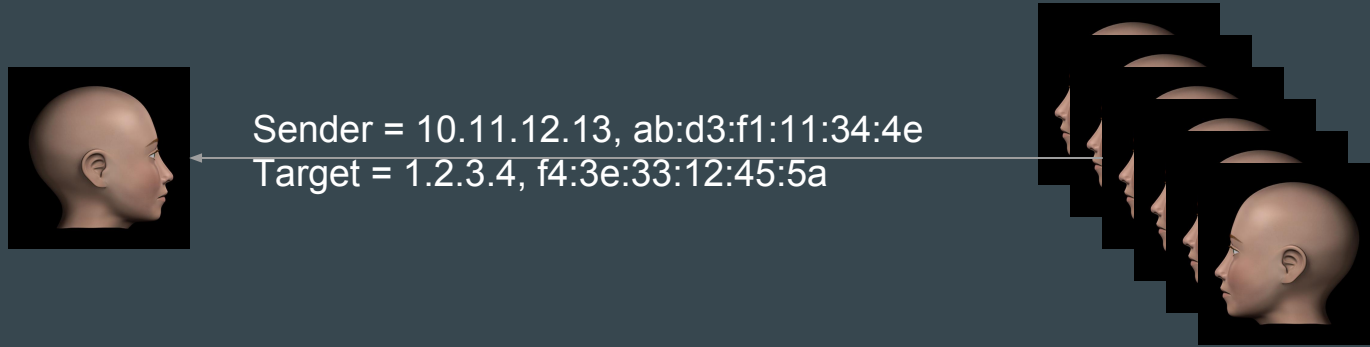
IP = 0x0800

0's in request

ARP



ARP



Sender address in Ethernet header does not have to be the same as the “sender’s HW address” in ARP response.

Possible way to detect an eavesdropper

Destination IP address = eavesdropper

Destination Ethernet address \neq eavesdropper

IP forwarding

Destination	Mask	Next hop	Interface
0.0.0.0	0.0.0.0	10.0.0.1	10.0.0.100
10.0.0.0	255.255.255.128	10.0.0.100	10.0.0.100

IP forwarding - algorithm

Let D be destination IP from packet header.

Let d_j be destination in j^{th} entry of forwarding table.

Let m_j be mask in j^{th} entry of forwarding table.

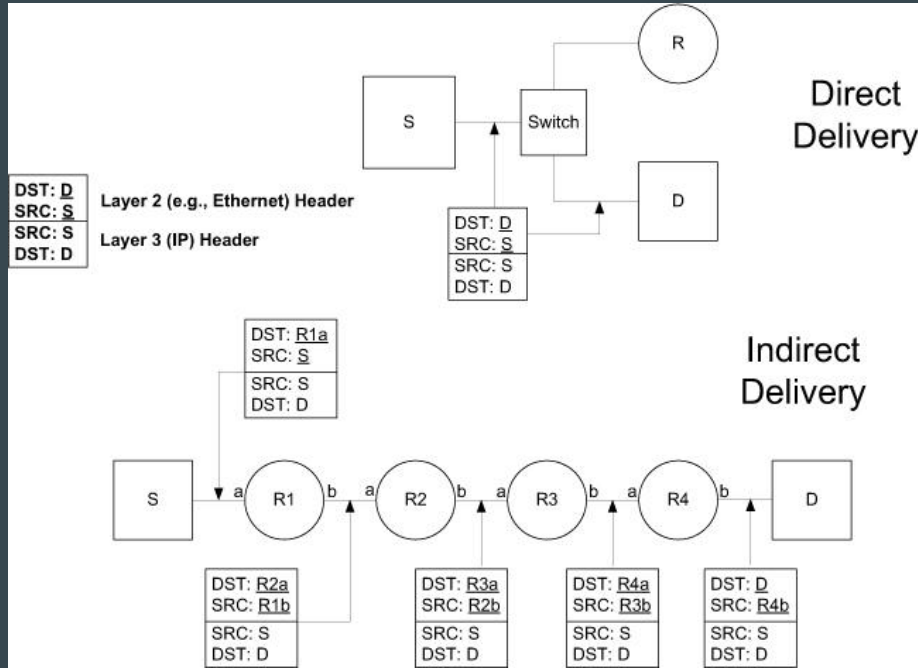
- Find all entries with: $(D \& m_j) = d_j$
- From amongst those, pick entry with most 1's in m_j
- > 1 best match \Rightarrow use some tie-breaking rule
- 0 matches \Rightarrow "host unreachable."

IP forwarding - examples

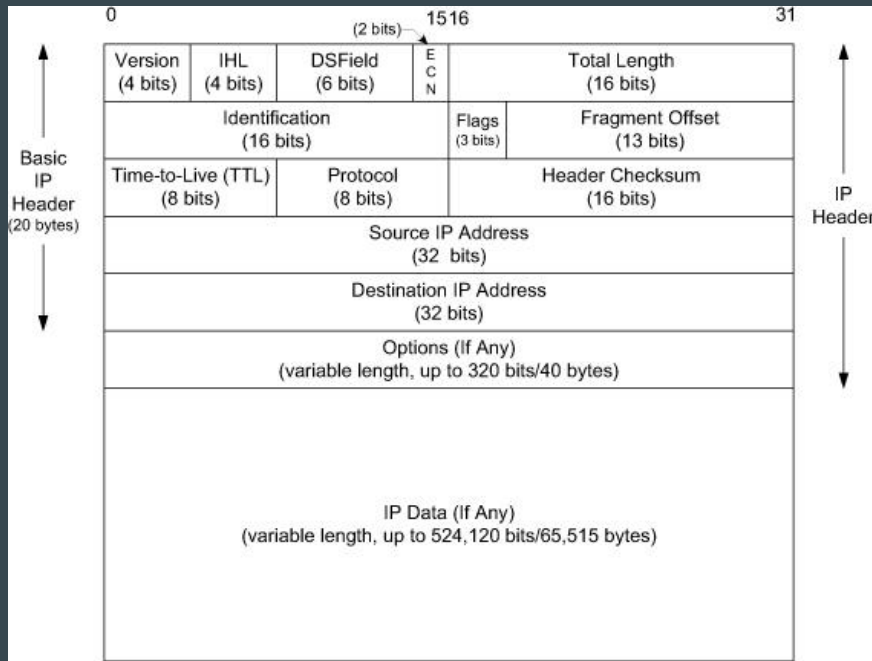
Destination	Mask	Next hop	Interface
0.0.0.0	0.0.0.0	10.0.0.1	10.0.0.100
10.0.0.0	255.255.255.128	10.0.0.100	10.0.0.100

- 10.0.0.19 matches both entries. 2nd entry is best-match.
- 178.162.3.4 matches 1st entry, not 2nd
- 10.0.0.131 matches 1st entry, not 2nd

IP forwarding - “direct” vs. “indirect” delivery




IP packet format



Fragmentation - the IP Header

All fragments of a datagram have the Same ID.



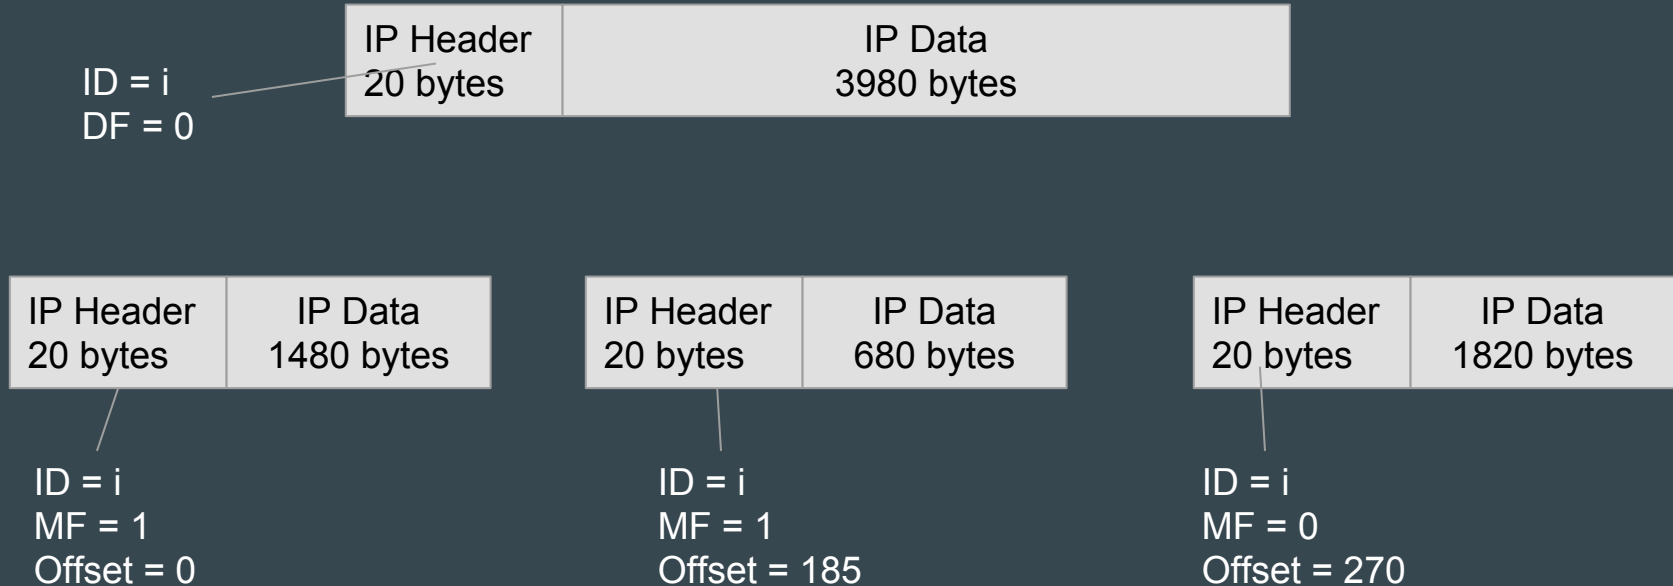
Reserved

Do not
Fragment

More
Fragments

Offset from start of data.
8-byte units.

Fragmentation - how it works via example



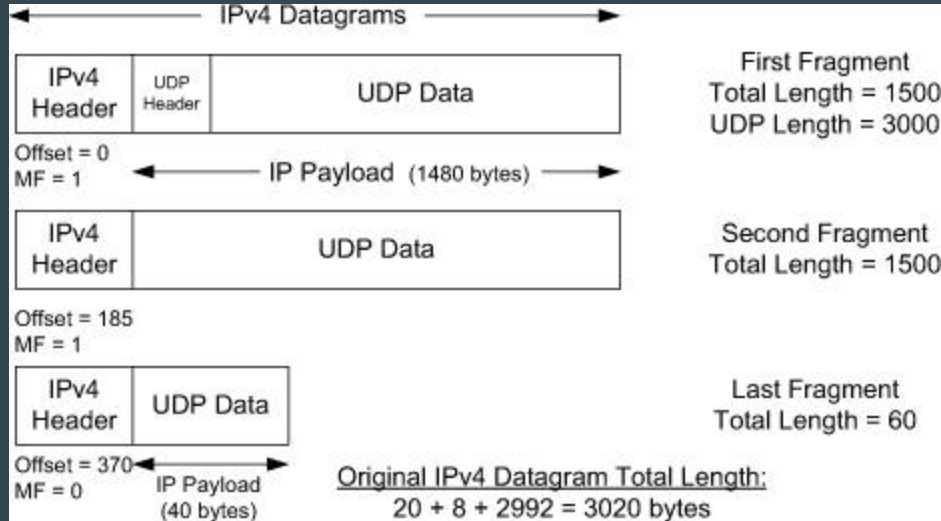
Fragmentation - issues (1)

- Complexity of algorithms \Rightarrow lots of bugs

See [RFC 791](#)

Fragmentation - issues (2)

- Higher layer protocol header is separated from data.



Fragmentation - issues (3)

- Identification field may not be long enough at 16 bits.
 - ID does not repeat for 120 seconds @ 1500 byte packets \Rightarrow 6.5 Mbps throughput
 - Throughput 1 Gbps @ 1500 byte packets \Rightarrow ID space exhausted in < 1 second
 - Proposed solution [RFC 6864](#): receiver relies on ID field for fragments only

Header Checksum

- “1’s complement of 1’s complement 16-bit sum.”

Checksum \leftarrow 00 00

Perceive header as chunks of 16 bits = 2 bytes

Checksum \leftarrow Checksum + (each of those 2 bytes)

Add carry back into Checksum

Checksum \leftarrow bit-complement of Checksum

Header Checksum, example

- E3 4F 23 96 44 27 99 F3

00 00
+ E3 4F
= E3 4F

E3 4F
+ 23 96
= 1 06 E5

1 06 E5
+ 44 27
= 1 4B 0C

1 4B 0C
+ 99 F3
= 1 E4 FF

E4 FF
+ 00 01
= E5 00

Checksum = 1A FF



Checking the checksum

- Compute checksum of entire header (including checksum)
- Result should be 00 00
- E.g.,
 - $(E3\ 4F) + (23\ 96) + (44\ 27) + (99\ F3) + (1A\ FF) = 1\ FF\ FE$
 - 1's complement: $(FF\ FE) + (00\ 01) = FF\ FF$
 - Complement = 00 00


TTL - Time To Live

- Use to limit datagram lifetime in the network
- Decrement by 1 @ each hop
- \Rightarrow header checksum must be recomputed @ every hop

Traceroute - a nifty use of TTL

- Objective - want to find path from here (source) to destination
- Method - leverage TTL field

Traceroute - how it works, via example



Traceroute - how it works, via example, Step 1



Traceroute - how it works, via example, Step 1



Source IP = Hop 1

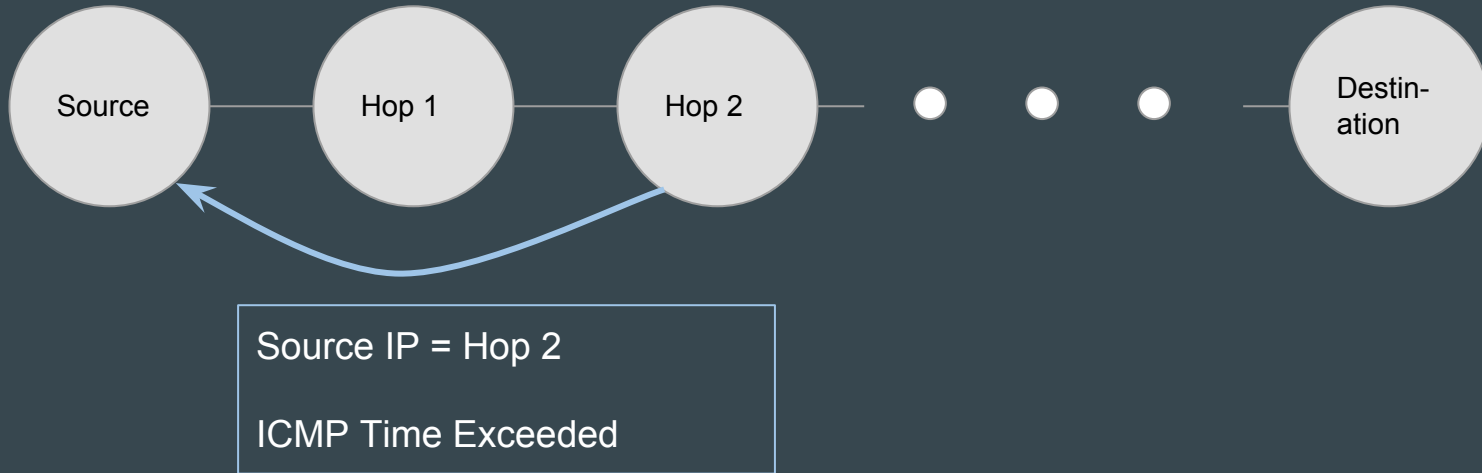
ICMP Time Exceeded
IP header + 8 bytes of data
from original datagram

Traceroute - how it works, via example, Step 1



Source IP = Source
Destination IP = Destination
TTL = 2

Traceroute - how it works, via example, Step 1



Question

It so happens that an IP packet sent by S to D incurs no errors in transit.

- True or false:

The value of the header checksum that D sees is the value that S put in the checksum field.