

# Interrupt Behavior

## IN THIS CHAPTER

Interrupt/Exception Sequences .....	145
Exception Exits .....	147
Nested Interrupts .....	148
Tail-Chaining Interrupts .....	148
Late Arrivals .....	149
More on the Exception Return Value .....	149
Interrupt Latency .....	152
Faults Related to Interrupts .....	152

## 9.1 INTERRUPT/EXCEPTION SEQUENCES

When an exception takes place, a number of things happen, such as

- Stacking (pushing eight registers' contents to stack)
- Vector fetch (reading the exception handler starting address from the vector table)
- Update of the stack pointer, link register (LR), and program counter (PC)

### 9.1.1 Stacking

When an exception takes place, the registers R0–R3, R12, LR, PC, and Program Status (PSR) are pushed to the stack. If the code that is running uses the Process Stack Pointer (PSP), the process stack will be used; if the code that is running uses the Main Stack Pointer (MSP), the main stack will be used. Afterward, the main stack will always be used during the handler, so all nested interrupts will use the main stack.

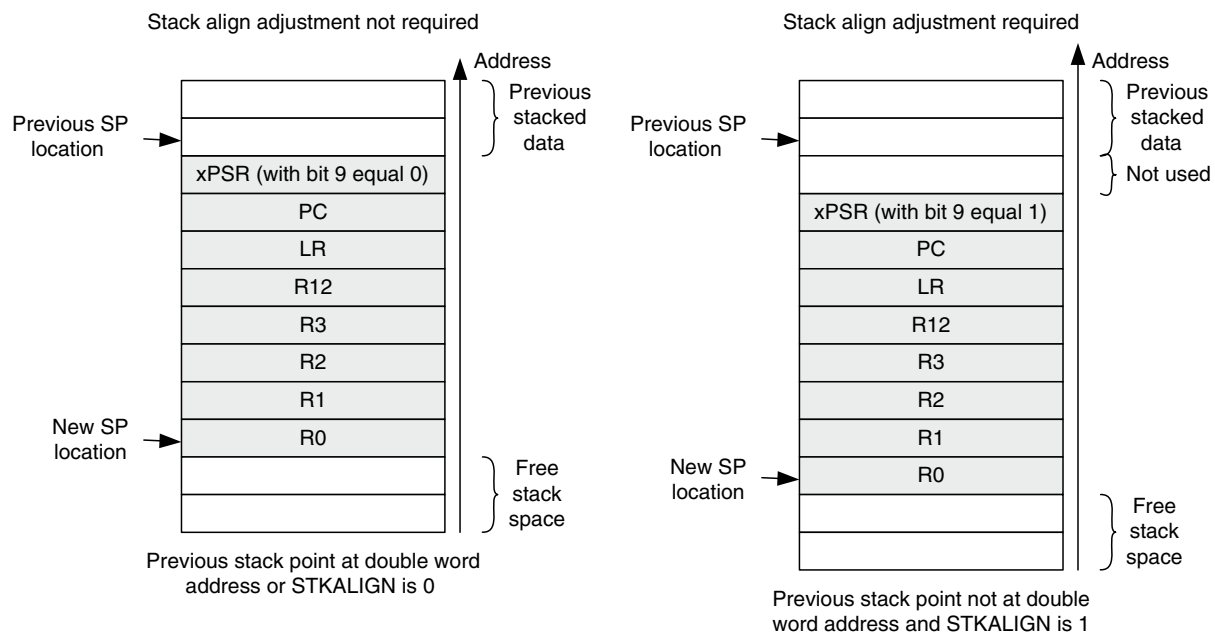
The block of eight words of data being pushed to the stack is commonly called a stack frame. Prior to Cortex™-M3 revision 2, the stack frame was started in any word address by default. In Cortex-M3 revision 2, the stack frame is aligned to double word address by default, although the alignment feature can be turned off by programming the STKALIGN bit in Nested Vectored Interrupt Controller (NVIC)

Configuration Control register to zero. The stack frame feature is also available in Cortex-M3 revision 1, but it needs to be enabled by writing 1 to the STKALIGN bit. More details on this register can be found in Chapter 12.

The data arrangement inside an exception stack frame is shown in Figure 9.1. The order of stacking is shown in Figure 9.2 (assuming that the stack pointer [SP] value is  $N$  after the exception). Due to the pipeline nature of the Advanced High-Performance Bus (AHB) interface, the address and data are offset by one pipeline state.


The values of PC and PSR are stacked first so that instruction fetch can be started early (which requires modification of PC) and the Interrupt Program Status register (IPSR) can be updated early. After stacking, SP will be updated, and the stacked data arrangement in the stack memory will look like Figure 9.1.

The reason the registers R0–R3, R12, LR, PC, and PSR are stacked is that these are caller-saved registers, according to C standards (C/C++ standard *Procedure Call Standard for the ARM Architecture*,



**FIGURE 9.1**

Exception Stack Frame.



**FIGURE 9.2**

Stacking Sequence.

AAPCS [Ref. 5]). This arrangement allows the interrupt handler to be a normal C function because registers that could be changed by the exception handler are saved in the stack.

The general registers (R0–R3 and R12) are located at the end of the stack frame so that they can be easily accessed using SP-related addressing. As a result, it's easy to pass parameters to software interrupts using stacked registers.

### 9.1.2 Vector Fetches

Although the data bus is busy stacking the registers, the instruction bus carries out another important task of the interrupt sequence: It fetches the exception vector (the starting address of the exception handler) from the vector table. Since the stacking and vector fetch are performed on separate bus interfaces, they can be carried out at the same time.

### 9.1.3 Register Updates

After the stacking and vector fetch are completed, the exception vector will start to execute. On entry of the exception handler, a number of registers will be updated. They are as follows:

- *SP*: The SP (either the MSP or the PSP) will be updated to the new location during stacking. During execution of the interrupt service routine, the MSP will be used if the stack is accessed.
- *PSR*: The IPSR (the lowest part of the PSR) will be updated to the new exception number.
- *PC*: This will change to the vector handler as the vector fetch completes and starts fetching instructions from the exception vector.
- *LR*: The LR will be updated to a special value called EXC\_RETURN.<sup>1</sup> This special value drives the interrupt return operation. The last 4 bits of the LR is used to provide exception return information. This is covered later in this chapter.

A number of other NVIC registers will also be updated. For example, the pending status of the exception will be cleared and the active bit of the exception will be set.

---

## 9.2 EXCEPTION EXITS

At the end of the exception handler, an exception exit (known as an *interrupt return* in some processors) is required to restore the system status so that the interrupted program can resume normal execution. There are three ways to trigger the interrupt return sequence; all of them use the special value stored in the LR in the beginning of the handler (see Table 9.1).

Some microprocessor architectures use special instructions for interrupt returns (for example, *reti* in 8051). In the Cortex-M3, a normal return instruction is used so that the whole interrupt handler can be implemented as a C subroutine.

When the interrupt return instruction is executed, the unstacking and the NVIC registers update processes that are listed in Table 9.1 are carried out.

---

<sup>1</sup>EXC\_RETURN has bit 31 to 4 all set to one (i.e., 0xfffffff). The last 4 bits define the return information. More information on the EXC\_RETURN value is covered later in this chapter.

**Table 9.1** Instructions That Can Be Used for Triggering Exception Return

Return Instruction	Description
<code>BX reg</code>	If the <code>EXC_RETURN</code> value is still in <code>LR</code> , we can use the <code>BX LR</code> instruction to perform the interrupt return.
<code>POP {PC}</code> , or <code>POP {..., PC}</code>	Very often the value of <code>LR</code> is pushed to the stack after entering the exception handler. We can use the <code>POP</code> instruction, either a single <code>POP</code> or multiple <code>POPs</code> , to put the <code>EXC_RETURN</code> value to the program counter. This will cause the processor to perform the interrupt return.
Load ( <code>LDR</code> ) or Load multiple ( <code>LDM</code> )	It is possible to produce an interrupt return using the <code>LDR</code> or <code>LDM</code> instruction with <code>PC</code> as the destination register.

1. *Unstacking*: The registers pushed to the stack will be restored. The order of the `POP` will be the same as in stacking. The stack pointer will also be changed back.
2. *NVIC register update*: The active bit of the exception will be cleared. For external interrupts, if the interrupt input is still asserted, the pending bit will be set again, causing it to reenter the interrupt handler.

### 9.3 NESTED INTERRUPTS

Nested interrupt support is built into the Cortex-M3 processor core and the NVIC. There is no need to use assembler wrapper code to enable nested interrupts. In fact, you do not have to do anything apart from setting up the appropriate priority level for each interrupt source. First, the NVIC in the Cortex-M3 processor sorts out the priority decoding for you. So when the processor is handling an exception, all other exceptions with the same or lower priority will be blocked. Second, the automatic hardware stacking and unstacking allow the nested interrupt handler to execute without risk of losing data in registers.

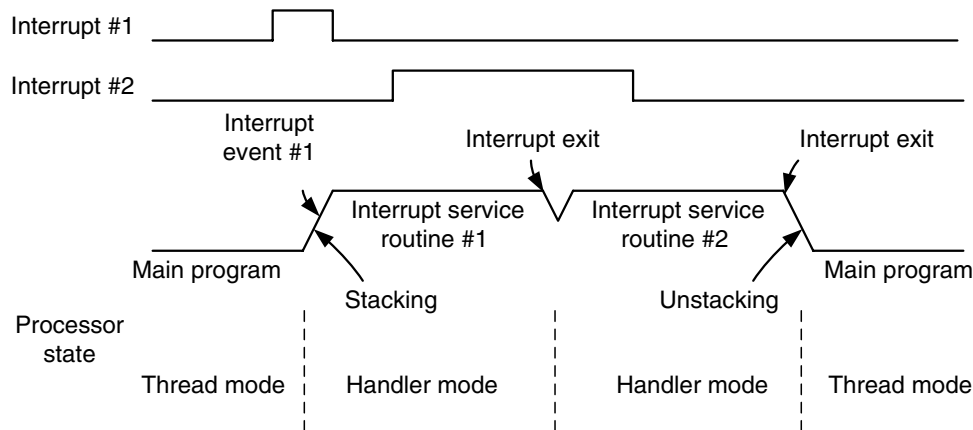
However, one thing needs to be taken care of: Make sure that there is enough space in the main stack if several nested interrupts are allowed. Since each exception level will use eight words of stack space and the exception handler code might require extra stack space, it might end up using more stack memory than expected.

Reentrant exceptions are not allowed in the Cortex-M3. Since each exception has a priority level assigned and, during exception handling, exceptions with the same or lower priority will be blocked, the same exception cannot be carried out until the handler is ended. For this reason, Supervisor Call (`SVC`) instructions cannot be used inside an `SVC` handler, since doing so will cause a fault exception.

### 9.4 TAIL-CHAINING INTERRUPTS

The Cortex-M3 uses a number of methods to improve interrupt latency. The first one we'll look at is *tail chaining* (see Figure 9.3).

When an exception takes place but the processor is handling another exception of the same or higher priority, the exception will enter pending state. When the processor has finished executing the current

**FIGURE 9.3**

Tail Chaining of Exceptions.

exception handler, it can then process the pended interrupt. Instead of restoring the registers back from the stack (unstacking) and then pushing them onto the stack again (stacking), the processor skips the unstacking and stacking steps and enters the exception handler of the pended exception as soon as possible. In this way, the timing gap between the two exception handlers is considerably reduced.

## 9.5 LATE ARRIVALS

Another feature that improves interrupt performance is *late arrival* exception handling. When an exception takes place and the processor has started the stacking process, and if during this delay a new exception arrives with higher preemption priority, the late arrival exception will be processed first.

For example, if Exception #1 (lower priority) takes place a few cycles before Exception #2 (higher priority), the processor will behave as shown in Figure 9.4, such that Handler #2 is executed as soon as the stacking completes.

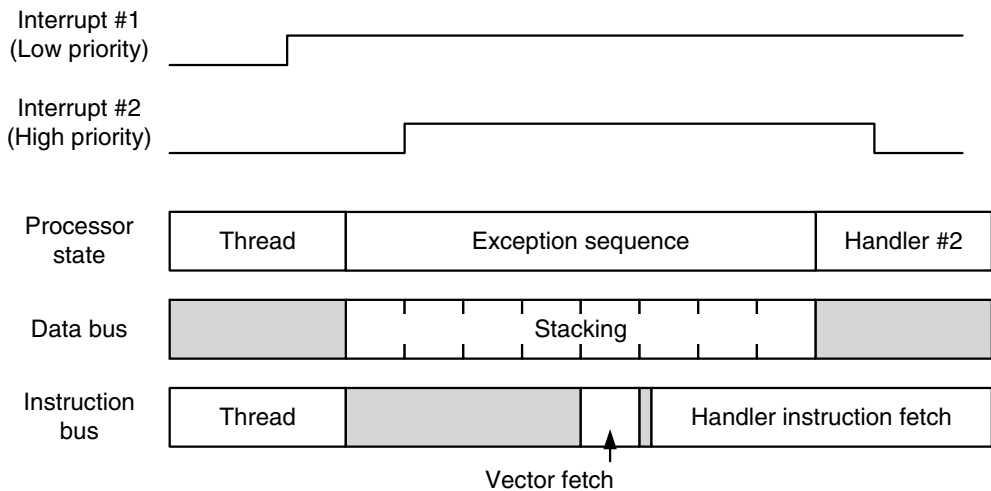
## 9.6 MORE ON THE EXCEPTION RETURN VALUE

When entering an exception handler, the LR is updated to a special value called EXC\_RETURN, with the upper 28 bits all set to 1. This value, when loaded into the PC at the end of the exception handler execution, will cause the processor to perform an exception return sequence.

The instructions that can be used to generate exception returns are as follows:

- POP/LDM
- LDR with PC as a destination
- BX with any register

The EXC\_RETURN value has bit (31:4) all set to 1, and bit (3:0) provides information required by the exception return operation (see Table 9.2). When the exception handler is entered, the LR value is updated automatically, so there is no need to generate these values manually.



**FIGURE 9.4**  
Late Arrival Exception Behavior.

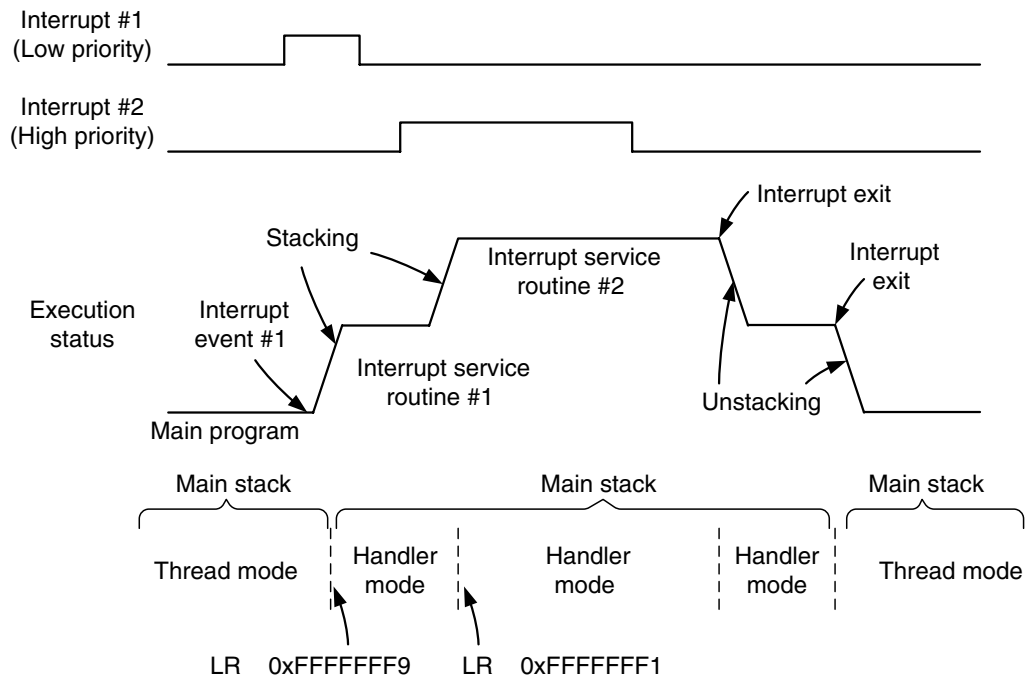
Table 9.2 Description of Bit Fields in EXC_RETURN Value					
Bits	31:4	3	2	1	0
Descriptions	0xFFFFFFFF	Return mode (thread/handler)	Return stack	Reserved; must be 0	Process state (Thumb/ARM)

Table 9.3 Allowed EXC_RETURN Values on Cortex-M3	
Value	Condition
0xFFFFFFFF1	Return to handler mode
0xFFFFFFFF9	Return to thread mode and on return use the main stack
0xFFFFFFFDD	Return to thread mode and on return use the process stack

Bit 0 indicates the process state being used after the exception return. Since the Cortex-M3 supports only the Thumb® state, bit 0 must be 1. The valid values (for the Cortex-M3) are shown in Table 9.3.

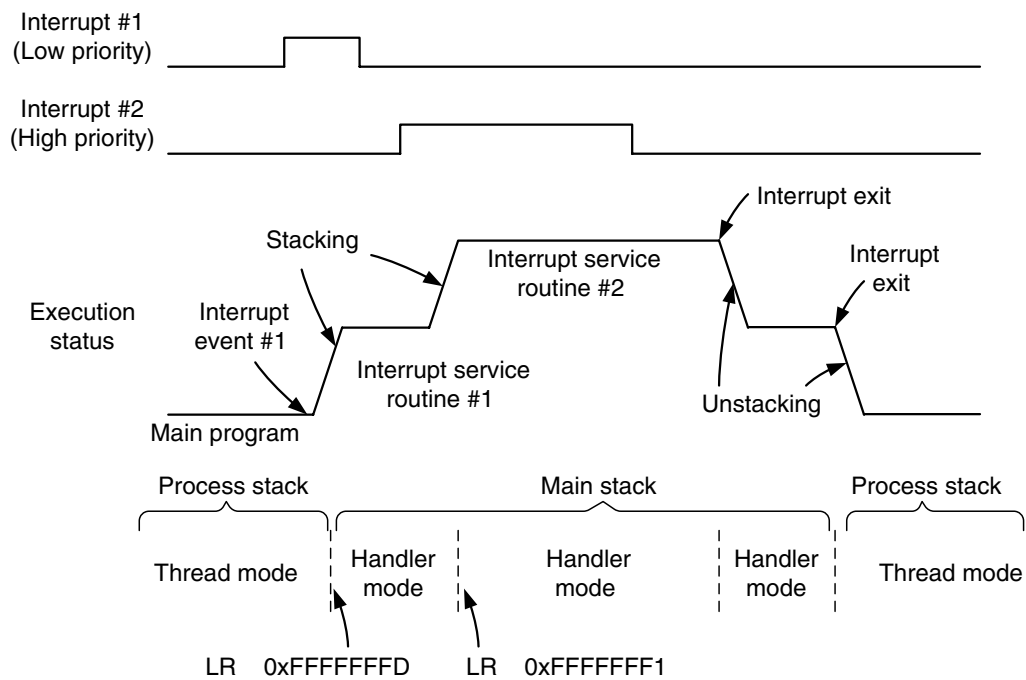
If the thread is using the MSP (main stack), the value of LR will be set to 0xFFFFFFFF9 when it enters an exception, and 0xFFFFFFFF1 when a nested exception is entered, as shown in Figure 9.5. If the thread is using PSP (process stack), the value of LR would be 0xFFFFFFFDD when entering the first exception and 0xFFFFFFFF1 for entering a nested exception, as shown in Figure 9.6.

As a result of the EXC\_RETURN number format, you cannot perform interrupt returns to an address in the 0xFFFFFFFF0–0xFFFFFFFFF memory range. However, since this address is in a nonexecutable region anyway, it is not a problem.



**FIGURE 9.5**

LR Set to EXC\_RETURN at Exception (Main Stack Used in Thread Mode).



**FIGURE 9.6**

LR Set to EXC\_RETURN at Exception (Process Stack Used in Thread Mode).

---

## 9.7 INTERRUPT LATENCY

The term *interrupt latency* refers to the delay from the start of the interrupt request to the start of interrupt handler execution. In the Cortex-M3 processor, if the memory system has zero latency, and provided that the bus system design allows vector fetch and stacking to happen at the same time, the interrupt latency can be as low as 12 cycles. This includes stacking the registers, vector fetch, and fetching instructions for the interrupt handler. However, this depends on memory access wait states and a few other factors.

For tail-chaining interrupts, since there is no need to carry out stacking operations, the latency of switching from one exception handler to another exception handler can be as low as six cycles.

When the processor is executing a multicycle instruction, such as divide, the instruction could be abandoned and restarted after the interrupt handler completes. This also applies to load double (LDRD) and store double (STRD) instructions.

To reduce exception latency, the Cortex-M3 processor allows exceptions in the middle of Multiple Load and Store instructions (LDM/STM). If the LDM/STM instruction is executing, the current memory accesses will be completed, and the next register number will be saved in the stacked xPSR (Interrupt-Continuable Instruction [ICI] bits). After the exception handler completes, the multiple load/store will resume from the point at which the transfer stopped. There is a corner case: If the multiple load/store instruction being interrupted is part of an IF-THEN (IT) instruction block, the load/store instruction will be cancelled and restarted when the interrupt is completed. This is because the ICI bits and IT execution status bits share the same space in the Execution Program Status Register (EPSR).

In addition, if there is an outstanding transfer on the bus interface, such as a buffered write, the processor will wait until the transfer is completed. This is necessary to ensure that a bus fault handler preempts the correct process.

Of course, the interrupt could be blocked if the processor is already executing another exception handler of the same or higher priority or if the Interrupt Mask register was masking the interrupt request. In these cases, the interrupt will be pended and will not be processed until the blocking is removed.

---

## 9.8 FAULTS RELATED TO INTERRUPTS

Various faults can be caused by exception handling. Let's take a look at these now.

### 9.8.1 Stacking

If a bus fault takes place during stacking, the stacking sequence will be terminated and the bus fault exception will be triggered or pended. If the bus fault is disabled, the hard fault handler will be executed. Otherwise, if the bus fault handler has higher priority than the original exception, the bus fault handler will be executed; if not, it will be pended until the original exception is completed. This scenario, called a *stacking error*, is indicated by the STKERR (bit 4) in the Bus Fault Status register (0xE000ED29).

If the stacking error is caused by a Memory Protection Unit (MPU) violation, the memory management fault handler will be executed and the MSTKERR (bit 4) in the Memory Management Fault Status register (0xE000ED28) will be set to indicate the problem. If the memory management fault is disabled, the hard fault handler will be executed.



### 9.8.2 Unstacking

If a bus fault takes place during unstacking (an interrupt return), the unstacking sequence will be terminated and the bus fault exception will be triggered or pended. If the bus fault is disabled, the hard fault handler will be executed. Otherwise, if the bus fault handler has higher priority than the current priority of the executing task (the core could already be executing another exception in a nested interrupt case), the bus fault handler will be executed. This scenario, called an *unstacking error*, is indicated by the UNSTKERR (bit 3) in the Bus Fault Status register (0xE000ED29).

Similarly, if the stacking error is caused by an MPU violation, the memory management fault handler will be executed and the MUNSTKERR (bit 3) in the Memory Management Fault Status register (0xE000ED28) will be set to indicate the problem. If the memory management fault is disabled, the hard fault handler will be executed.

### 9.8.3 Vector Fetches

If a bus fault or memory management fault takes place during a vector fetch, the hard fault handler will be executed. This is indicated by VECTTBL (bit 1) in the Hard Fault Status register (0xE000ED2C).

### 9.8.4 Invalid Returns

If the EXC\_RETURN number is invalid or does not match the state of the processor (as in using 0xFFFFFFFF1 to return to thread mode), it will trigger the usage fault. If the usage fault handler is not enabled, the hard fault handler will be executed instead. The INVPC bit (bit 2) or INVSTATE (bit 1) bit in the Usage Fault Status register (0xE000ED2A) will be set, depending on the actual cause of the fault.

## 9: Interrupt Behavior

### 9.1: Interrupt/Exception

#### 9.1.1: Stacking

When an exception takes place registers are pushed to the stack

- R0-R3
- R12
- LR (Link Register)
- PC (Program Counter)
- PSR (Program Status Register)

The stack to which these will be pushed will depend on the stack pointer that the program was using (Process Stack Pointer goes to process stack and Main Stack Pointer goes to main stack). Interrupts always use the main stack.

The block of eight words (the eight registers pushed at the start of an interrupt) is called a stack frame.

We stack the PC and PSR are stacked first so that we can start the instruction fetch early for efficiency. This also allows us to update the Interrupt Program Status Register early. Once we finish stacking all of the new data the Program Stack is updated to its new position at the top of the new data.

Registers R0-R3, R12, LR, PC, PSR are all stacked to abide by C standards so that the interrupt handler can be a C function (any registers that would be changed by the exception handler are already saved).

#### 9.1.2: Vector Fetches

While the data bus is stacking registers the instruction bus fetches the exception vector from the vector table. Since these are done on separate buses they can be done concurrently.

#### 9.1.3: Register Updates

When we enter the exception handler we update registers:

- the SP (either main stack pointer or program stack pointer) is moved the new location after stacking all necessary registers
- the PSR's lowest part the IPSR is updated to the new exception number
- PC changes to the vector handler and starts fetching exception instructions
- LR is updated to a value EXC\_RETURN which drives the interrupt operation, its last 4 bits provide exception return information

## 9.2: Exception Exits

When the exception handler is done an exit (or interrupt return) is required to return to the regular program. There are three ways to trigger this all of which use the special value stored in LR:

- If LR still has EXC\_RETURN in it we can use BX LR to link back the the program
- If the LR has been pushed to the stack after entering the handler we can use a series of pops to put it on the PC
- We can use a LDR to load the correct address into the PC

After this the registers are unstacked (pop in the same order as pushing) and the NVIC register is updated (this checks if the interrupt is done running if it is not cleared the program will reenter).

## 9.3: Nested Interrupt

The ability to support nested interrupts is build into the processor core that we will be using, we don't need to do anything apart from setting up the appropriate priority level for each interrupt source. While an exception is running all ones of lower priority will be ignored. Since the hardware automatically stacks and unstacks for us we can do this with no risk of data loss.

We do need to make sure that there is enough space in the main stack for several nested interrupts. This also means that we cannot reenter an exception since the priority level would be the same and this make it ignored.

## 9.4: Tail-Chaining Interrupts

This is a method to improve interrupt latency. If an exception takes place but the processor is busy with one of a higher priority it enters a pending state. This won't be processed until the processor is done with its current exception. When it does, the processor doesn't bother unstacking and restacking the registers since the program data hast changed.

## 9.5: Late arrivals

If exceptions happen during the stacking process and it is interrupted by one with a higher priority the processor will jump straight to that one instead of starting the lower priority one and then going to the higher one.

## 9.6: More on the Exception Return Value

When we enter an exception we set the LR to EXC\_RETURN (upper 28 bits set to 1). This value is then loaded into the PC at the end of the handler which tells the processor to execute the return sequence. This can be triggered by:

- POP/LDM from the stack
- LDR the PC

- BX with a register

The EXC\_RETURN's bottom 4 bits provide information required by the operation used to execute it (this updates the LR manually so we don't have to generate these values). The last bit of EXC\_RETURN indicates the process state being used after the exception returns.

- 0xfffff1: return to handler mode (if a nested exception was entered)
- 0xfffff9: return to thread mode and on return use the main stack (if the thread was using MSP)
- 0xfffffD: return to thread mode and on return use the process stack (if the thread was using PSP)

This means that you cannot have a interrupt return to an address in the range 0xfffff0-0xfffff, but these addresses are in a non-executable region anyway.

## 9.7: Interrupt Latency

Latency is the delay from the start of an interrupt request to start of handler execution. The lowest possible latency for our processor is 12 cycles (if the memory system has no latency, and the bus system lets us fetch the vector while the memory bus does stacking). If we are executing tail-chaining interrupts (where stacking isn't required) this can be as low as 6 cycles. For instructions that are multicycle we could abandon it and restart it when the handler is done. The processor allows exceptions during multiple load and multiple store operations, the current memory access is completed and the next register number is stored in the xPSR, once it's done the operation picks up where it left off. If this happens in a IF/THEN, the LDM/STM is killed and restarted, this is because the if/then status is saved in the same place as the ICS (the next memory to be accessed by the operation).

When there is an outstanding transfer on the bus interface it is allowed to finish to ensure that the bus fault handler overrides the current process.

## 9.8: Faults Related to Interrupts

Faults can be caused by exception handling.

### 9.8.1: Stacking

A bus fault can happen during stacking which will result in the stacking sequence dying and the bus fault being triggered. This is called a stacking error, indicated by the STKERR (4th bit in the Bus Fault Status Register 0xE000ED29).

If the stacking error is caused by a memory protection unit violation the memory management handler will be executed and the MSTKERR (bit 4 of the Memory Management Fault Status register 0xE000ED28).

### 9.8.2: Unstacking

If a bus fault takes place during unstacking the exact thing happens but you set the 3rd bits of the Bus Fault Status or Memory Management Fault Status register.

### **9.8.3: Vector Fetches**

If a bus fault or memory management fault takes place during a vector fetch the hard fault handler is executed, indicated by VECTTBL(bit i in Hard Fault Status register, 0xE000ED2C)

### **9.8.4: Invalid Returns**

If the EXC\_RETURN is invalid or doesn't match the state of the processor it triggers the usage fault, the INVPC bit 2 or INVSTATE bit 1 in the Usage Faulted Status register 0xE000ED2A is set.