



SE350: Operating Systems

Mahmoud Salem

m4salem@uwaterloo.ca

Please email through Learn

Office hours will be decided if needed

Tutorial

- Chapter 5 :
 - Review Questions: 5.1 – 5.7
 - Problems: 5.1, 5.3, 5.7, 5.8, 5.12, 5.21

7th and 8th Editions have same numbering for those problems.

Chapter 5 Concurrency: Mutual Exclusion and Synchronization

Reference:

Operating Systems “Internals and Design Principles” 8th Edition – William Stallings

Design issues raised by existence of concurrency – handled by OS

- Keep track of various processes
 - this is handled by process control blocks
- OS must allocate and de-allocate resources
 - processing time, memory, files, I/O, ..
- Protection of data and resources from unintended interference by other processes.
- Function of process and output produced must be independent of the speed relative to the speed of other concurrent processes.

Concurrency context

- Multiple applications [Multiprogramming]
 - processing time is dynamically shared among active applications
- Structure applications
 - some applications can be effectively programmed as a set of concurrent processes
- OS Structure
 - OS itself often implemented as set of processes and threads

Mutual Exclusion

- The basic requirement for execution of concurrent processes is to enforce mutual-exclusion.
- It's the requirement that when one process is in the critical section that access shared resources, no other process may be in a critical section that access any of those shared resources.

Degrees of awareness between processes

- Processes unaware of each other
 - independent, non intended to work together
 - OS regulate their **competition** for resources
- Processes indirectly aware of each other
 - not necessarily aware of each other w.r.t to IDs
 - but might share access to same object as I/O buffer, have to **cooperate**
- Processes directly aware of each other
 - communicate with each other (**cooperate**) by process ID to work jointly on some activity

Cooperation vs Competition

- Competing processes:
 - need access to the same resource at the same time as file or printer
- Cooperating processes:
 - Either share access to common object
 - Such as memory buffer
 - Or able to communicate with each other
 - performance of some application

Check Table 5.2 in text book !

Control problems with competing processes

- Mutual Exclusion
 - need access to same resource at same time
 - critical section code (as printing code) - critical resource (as a printer)
- Deadlock
 - competing process need exclusive access to more than one resource
 - situation where two or more processes are unable to proceed because each is waiting for one of the others
- Starvation
 - a competing process which is ready to proceed may be indefinitely denied access to needed resource because of other processes monopolizing this resource. (overlooked by the dispatcher)

Problems

Problem 5.1

- The same problems of concurrency are present for both multiprogramming and multiprocessing.
 - Mention two differences in terms of concurrency between multiprogramming and multiprocessing ?

Problem 5.1 Solution

- Multiprogramming
 - Interleaving of execution on uniprocessor
 - Concurrency is handled by disabling interrupts
- Multiprocessing
 - Overlapping of execution on multiple processor
 - Concurrency is handled by locking shared objects
 - Ex: Memory accessed by multiple processors at the same time.

Problem 5.3

```
P1: {  
shared int x;  
x = 10;  
while (1) {  
    x = x - 1;  
    x = x + 1;  
    if (x != 10)  
        printf("x is %d",x)  
    }  
}
```

```
P2: {  
shared int x;  
x = 10;  
while ( 1 ) {  
    x = x - 1;  
    x = x + 1;  
    if (x!=10)  
        printf("x is %d",x)  
    }  
}
```

- Consider this program running on a uniprocessor system, the scheduler can run the instructions in pseudo parallel fashion.
 - Show an execution trace to print “x is 10” ?

Solution 5.3a

- The key here is to find a scenario that reach the check before the print statement where $x \neq 10$, then by the time the print statement is executed the x should have been equal to 10.

P1	P2	Value of X
$X = X - 1$		9
$X = X + 1$		10
	$X = X - 1$	9
If($x \neq 10$)		9
	$X = X + 1$	10
print X		10
	If($x \neq 10$)	10

Problem 5.3 - continued

- Show a sequence such that the statement “x is 8” printed.

Knowing that the increment and decrement are not done atomically in the assembly language code as follows:

```
LD      R0,X    /* load R0 from memory location x */
INCR    R0      /* increment R0 */
STO     R0,X    /* store the incremented value back in X */
```

Solution 5.3b

- The key here is to find a scenario where you a the loop iteration decreases the value of X so eventually you will print the value of 8.

P1		P2		Values
X = X -1				V(X) = 9
		X = X-1		V(X) = 8
X= X+1	LD R01,X			V(RO1) = 8
	INC RO1			V(RO1) = 9
		X= X+1	LD R02,X	V(RO2) = 8
			INC RO2	V(RO2) = 9
			STO RO2,X	V(X) = 9
	STO RO1,X			V(X) = 9
Print X				V(X) = 9
		Print X		V(X) = 9

Repeat this scenario for more iteration to print "X is 8"

Problem 5.7

- Lamport's Bakery Algorithm

```
boolean choosing[n];
int number[n];
while (true) {
    choosing[i] = true;
    number[i] = 1 + getmax(number[], n);
    choosing[i] = false;
    for (int j = 0; j < n; j++) {
        while (choosing[j]) { };
        while ((number[j] != 0) && (number[j], j) < (number[i], i)) { };
    }
    /* critical section */;
    number[i] = 0;
    /* remainder */;
}
```

- *choosing*[] is initialized to false, and *number*[] is initialized to 0.
- The i^{th} element of an array is only writeable by process i , can be read by any.
- The notation $(a,b) < (c,d)$ defined as:
$$(a < c) \text{ or } (a = c \text{ and } b < d)$$

Problem 5.7 - continued

- Describe the algorithm in words ?
- Show that this algorithm avoids deadlock ?
- Show that it enforces mutual exclusion ?

5.7 Solution

- For a process i :
 - $\text{choosing}[i] = \text{true}$ means the process is taking a number
 - $\text{number}[i] = 1 + \max()$ means largest number assigned to process after increment.
 - $\text{choosing}[i] = \text{false}$ means the process took a number
 - Loop to check all processes:
 - No process is in procedure of taking a number
 - For the processes holding a number, wait until you are having the lowest number. If you have number equal to other process, the process with lower index proceed.
 - proceed with the critical section
 - process finished the critical section and reset its ticket

5.7 Solution - continued

- Deadlock

Assuming that the processes are assigned a unique index, no two processes can block on each other. Always we will have a process that proceed.

5.7 Solution - continued

- Mutual Exclusion

Let's say P1 in critical section, that means:

- P1 has found choosing[p2index] is false
 - Either P2 just entered main loop – didn't get number [OK]
 - Or P2 was already assigned a number [go to next check]
- P1 has found it got smaller number than P2, or if got the same number it already have lower index [OK]

P2 can't enter the critical section as P1 is the only satisfier to the conditions until it resets its number[p1index] to zero at end of critical section.

Problem 5.21

Any effect on program meaning if statements are as follows ?

- `semWait(s);`
`semWait(e);`
- `semSignal(n);`
`semSignal(s);`
- `semWait(s);`
`semWait(n);`
- `semSignal(e);`
`semSignal(s);`

```
/* program boundedbuffer */
const int sizeofbuffer = /* buffer size */;
semaphore s = 1;
semaphore n = 0;
semaphore e = sizeofbuffer;
void producer()
{
    while (true)
    {
        produce();
        semWait(e);
        semWait(s);
        append();
        semSignal(s);
        semSignal(n)
    }
}
void consumer()
{
    while (true)
    {
        semWait(n);
        semWait(s);
        take();
        semSignal(s);
        semSignal(e);
        consume();
    }
}
void main()
{
    parbegin (producer, consumer);
}
```

5.21 Solution

- a) Deadlock might occur when buffer is full ($e=0$).
- b) No issues, although optimally you should have the critical section with the main functionality (append/take).
- c) Deadlock as the producer won't reach the signal call when the buffer is empty ($n=0$)
- d) Same as b.

Additional Problems

Problem 5.8

```
1 int number[n];
2 while (true) {
3     number[i] = 1 + getmax(number[], n);
4     for (int j = 0; j < n; j++) {
5         while ((number[j] != 0) && (number[j], j) < (number[i], i)) { };
6     }
7     /* critical section */;
8     number[i] = 0;
9     /* remainder */;
10 }
```

- Modified version of Lamport Bakery algorithm without using *choosing variable*. Does this violate mutual exclusion ?

5.8 Solution

- The `getmax()` will have to read the `number[]` elements before assigning the incremented value to a process.
- Assume P0 and P1 both read the `max number[]` as 0 (initial value) at the same time
- P1 proceeds with value 1 and enter C.S
- P0 proceeds with value 1 and enter C.S because it has lower index and same number as P1.
- **Mutual Exclusion is violated.**
- The *choosing* variable with it's loop prevents a process from performing the C.S check while another process is being assigned a number.

Problem 5.12

- Any difference for using this semaphore definition compared to the one in Fig 5.3 ?

```
void semWait(s)
{
    if (s.count > 0) {
        s.count--;
    }
    else {
        place this process in s.queue;
        block;
    }
}

void semSignal (s)
{
    if (there is at least one process blocked on
        semaphore s) {
        remove a process P from s.queue;
        place process P on ready list;
    }
    else
        s.count++;
}
```

5.12 Solution

- The definition in Fig 5.3 allows negative values for the count variable, which counts the waiting processes. The provided definition doesn't have this info.
- Using this semaphore is the same as to the one in Fig 5.3 as they function in the same way.

Helping tips for understanding semaphores

- Understand requirements of semaphores in page 210 in 8th edition or page 209 in 7th edition.
- Understand the effect of counter initialization.
- Understand the effect of Signal/Wait order

Questions ?