

Chapter 10

Multiprocessor and Real-Time Scheduling


(based on original slides by Pearson)

Classifications of Multiprocessor Systems

- Loosely coupled or distributed multiprocessor, or cluster
 - Each processor has its own memory and I/O channels
 - Classical distributed system
- Functionally specialized processors
 - Such as I/O processor
 - Controlled by a master processor
 - FPGA cards

Classifications of Multiprocessor Systems

- Tightly coupled multiprocessing
 - Processors share main memory
 - Controlled by operating system
 - Multi-core, SMP, cell processor, FPGA in processor socket



Independent Parallelism

- Processes are separate applications
- → no synchronization among processes
- Example is time-sharing system
 - Run word processing, shell, mail, browser
 - Properties similar to the uniprocessor system

Types of Parallelism

- Coarse and Very Coarse-Grained Parallelism
 - Distributed processing across network nodes
 - Multiprocessing of concurrent processes in a multiprogramming environment
 - Little synchronization among processes
 - Good for concurrent processes running on a multiprogrammed uniprocessor
 - Can be supported on a multiprocessor with little change

Types of Parallelism

- Medium-Grained Parallelism
 - Parallel processing or multitasking within a single application
 - Threads usually interact frequently and share data
 - → requires synchronization
 - Scheduling decisions regarding one thread might affect the performance of the entire application.

Types of Parallelism

- Fine-Grained Parallelism
 - Parallelism inherent in a single instruction stream
 - Highly parallel applications

Scheduling Design Issues

- Scheduler needs to consider:
 - Assignment of processes to processors
 - Use of multiprogramming on individual processors
 - Actual dispatching of a process

Assignment of Processes to Processors

- Two architectural styles:
 1. Multiprocessor is uniform
 2. Multiprocessor is heterogeneous
- Assuming architectural style 1:
 - a) Assign processes to a dedicated processor
 - b) Migrate processes between processors
- Style 2 requires special software.

Assignment of Processes to Processors

- Option 1.a: static assignment
 - Low system overhead, because decision is made once.
 - Permits group or gang scheduling
 - Processor can be idle while another has a backlog
- Option 1.b: process migration
 - Potential high overhead due to migration
- Option 1.ab: dynamic load balancing
 - Have a static assignment, but migrate processes sometimes

Assignment of Processes to Processors

- Master/slave architecture
 - Key kernel functions always run on a particular processor (→ RTLinux design)
 - Master is responsible for scheduling
 - Slave sends service request to the master and waits for result
 - Advantage: simple design, similar to uniproc.
 - Disadvantages
 - Failure of master brings down whole system
 - Master can become a performance bottleneck

Assignment of Processes to Processors

- Peer architecture
 - Kernel can execute on any processor
 - Each processor does self-scheduling
 - Complicates the operating system
 - Make sure two processors do not choose the same process


Multiprogramming

- Do we still need multiprogramming in a multiprocessor environment?
- Consider an 80 core machine
- No definitive answer. Advantages and disadvantages exist.
 - Low system overhead, less complexity
- Btw. in-application concurrency still necessary

Process Dispatching


- Do we still require sophisticated concepts to guide the scheduling decisions?
 - Complex scheduling algorithms
 - Priorities
 - Feedback queues
 - Compute metrics at run time
- Some are unnecessary and even counter productive → active area of research

Comparison One and Two Processors



(a) Comparison of RR and FCFS

Comparison One and Two Processors



(b) Comparison of SRT and FCFS

Thread Scheduling

- Separate execution from resource ownership
- An application consists of multiple cooperating, concurrently-executing threads
- Uniprocessor:
 - Program structuring aid
 - Overlap I/O with processing
 - Low management overhead (compared to MP)
- Multiprocessor:
 - True parallelism

Thread Scheduling & Assignment Overview

- Load sharing
 - Processes are not assigned to a particular processor
 - Load sharing vs load balancing
- Gang scheduling
 - A set of related threads is scheduled to run on a set of processors at the same time

Thread Scheduling & Assignment Overview

- Dedicated processor assignment
 - Threads are assigned to a specific processor
- Dynamic scheduling
 - Number of threads can be altered during course of execution

Load Sharing

- Global queue, processor picks process
- Load is distributed evenly across the processors
- Advantages:
 - No processor is left idle while there are processes available.
 - No centralized scheduler required
 - Transparent for the developer

Disadvantages of Load Sharing



- Central queue needs mutual exclusion
- Preemptive threads are unlikely to resume execution on the same processor (cache misses)
- If all threads are in the global queue, all threads of a program will not gain access to the processors at the same time

Gang Scheduling

- Threads often need to synchronize with each other
- → run multiple threads together
- Simultaneous scheduling of threads that make up a single process
- Useful for applications where performance severely degrades when any part of the application is not running

Processor Allocation

- Gang scheduling can lead to inefficient use of the multiprocessor
 - Group independent threads into groups
 - Use weights in the scheduling algorithm



Dedicated Processor Assignment

- Dedicate thread groups to processors until application completes
- No multiprogramming !
- Looks bad: extremely wasteful
- However:
 - Assume 1000 processor cores
 - Zero overhead

Application Speedup




Figure 10.4 Application Speedup as a Function of Number of Threads

Dynamic Scheduling

- Number of threads in a process are altered dynamically by the application
- Requires a layer of indirection which maps computation tasks to threads

Real-Time Scheduling

- Correctness of the system depends on
 - Logic behavior
 - Timing
- **A correct value at the wrong time is a fault.**
- Tasks or processes attempt to control or react to events that take place in the outside world
- These events occur in “real time” and tasks must be able to keep up with them

Properties Of RT Systems

- Real-time is about producing the correct result at the right time.

| Value | Timing | Result |
|----------------|----------------|---------------|
| Wrong | Too late | Failure |
| Wrong | On time | Failure |
| Correct | Too late | Failure |
| Correct | On time | Ok |

- Temporal constraints are a way to specify, when the value is on time.

Soft Temporal Constraints

- A **soft real-time system** is one where the response time is normally specified as an average value. This time is normally dictated by the business or market.
- A single computation arriving late is not significant to the operation of the system, though many late arrivals might be.
- Ex: Airline reservation system - If a single computation is late, the system's response time may lag. However, the only consequence would be a frustrated potential passenger.

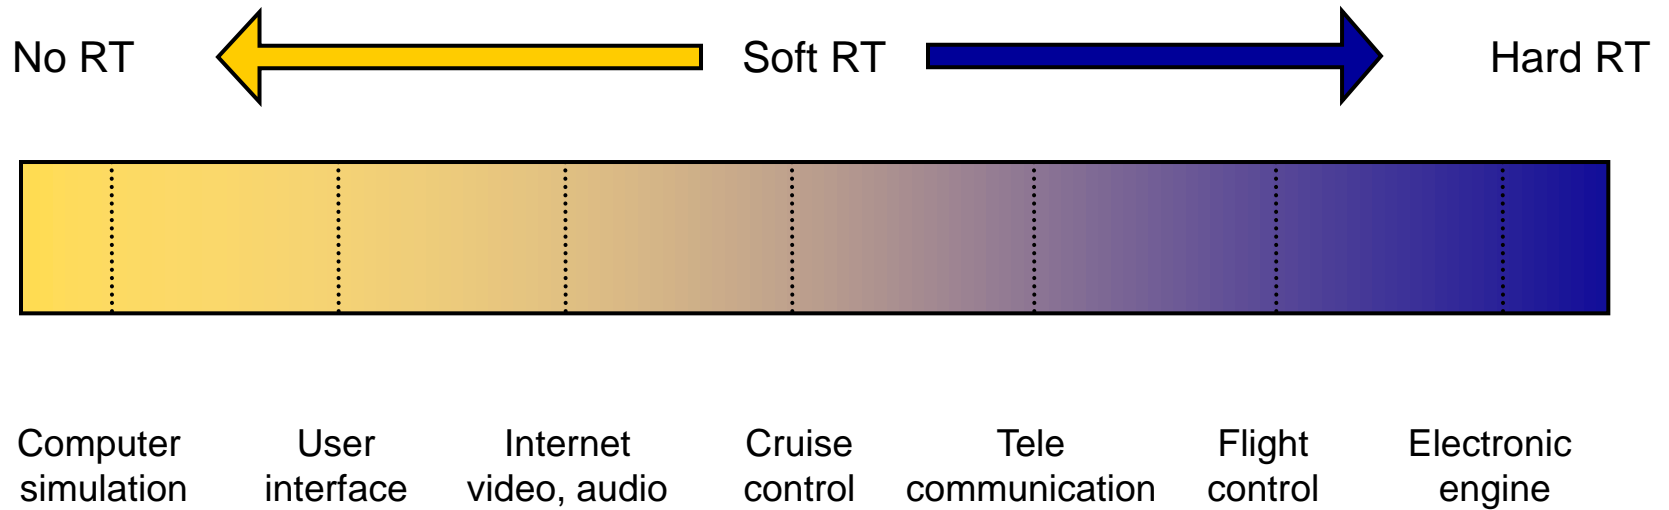
Hard Temporal Constraints

- A **hard real-time system** is one where the response time is specified as an absolute value. This time is normally dictated by the environment.
- A system is called a hard real-time if tasks always must finish execution before their deadlines or if message always can be delivered within a specified time interval.
- Hard real-time is often associated with safety critical applications. A failure (e.g. missing a deadline) in a safety-critical application can lead to loss of human life or severe economical damage.

Firm Temporal Constraints

- In a **firm real-time system** timing requirements are a combination of both hard and soft ones. Typically the computation will have a shorter soft requirement and a longer hard requirement.
- Ex: Ventilator – The system must ventilate a patient so many times within a given time period. But a few second delay in the initiation of the patient's breath is allowed, but not more.

Real-Time Spectrum



Characteristics

- Determinism
 - Operations are performed at **fixed, predetermined times** or within **predetermined time intervals**
 - Example: Concerned with how long the operating system delays before acknowledging an interrupt and there is sufficient capacity to handle all the requests within the required time

Characteristics

- Responsiveness of tasks/jobs
 - How long, after acknowledgment, it takes the operating system to service the interrupt
 - Includes **amount of time to begin execution** of the interrupt
 - Includes **the amount of time to perform** the interrupt service routine
 - Includes **interference** from interrupt nesting

Characteristics

- User control
 - Usually user has **little control** (priorities, grouping), in an RTOS this is different
 - User specifies priority, importance, timing
 - Manually control memory management (locking pages, specifying resource demands)
 - Manually controlling I/O algorithms (What disk transfer algorithms to use)

Characteristics

- Reliability
 - Degradation of performance may have catastrophic consequences (safety-critical systems)
- Fail-soft operation
 - Ability of a system to fail in such a way as to preserve as much capability and data as possible
 - Graceful degradation
 - Compare to fail-hard, fail-safe, fail-silent

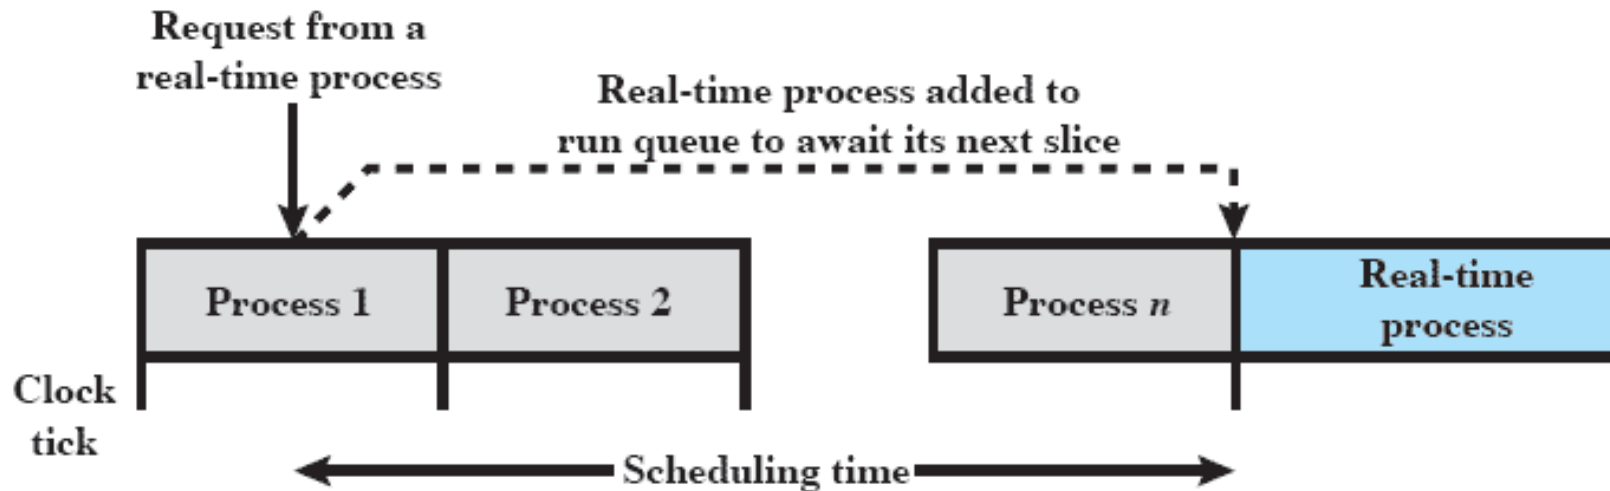
Features of Real-Time OS

- Fast process or thread switch
- Small size
- Ability to respond to external interrupts quickly
- Multitasking with interprocess communication tools such as semaphores, signals, and events

Features of Real-Time OS

- Use of special sequential files that can accumulate data at a fast rate
- Preemptive scheduling based on priority
- Minimization of intervals during which interrupts are disabled
- Delay tasks for fixed amount of time
- Special alarms and timeouts

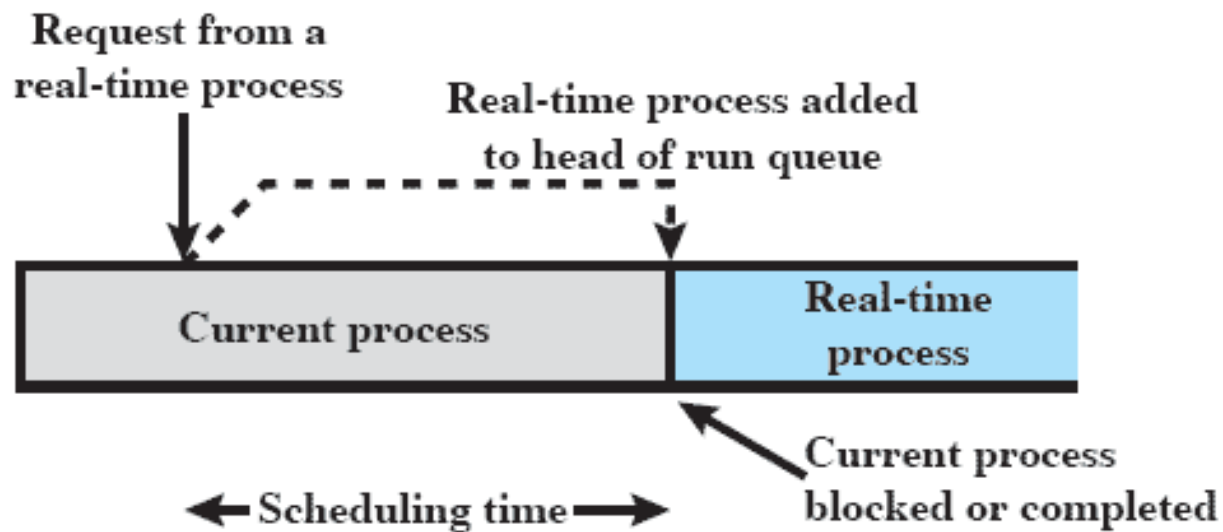
Scheduling of Real-Time Process



(a) Round-robin Preemptive Scheduler

Unacceptable!

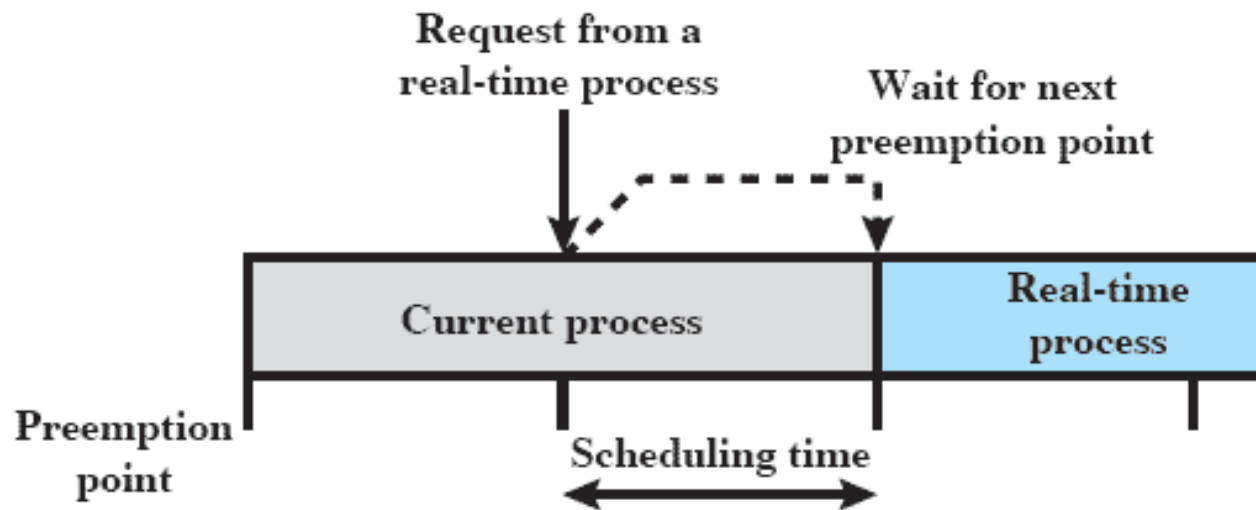
Scheduling of Real-Time Process



(b) Priority-Driven Nonpreemptive Scheduler

Unacceptable too!

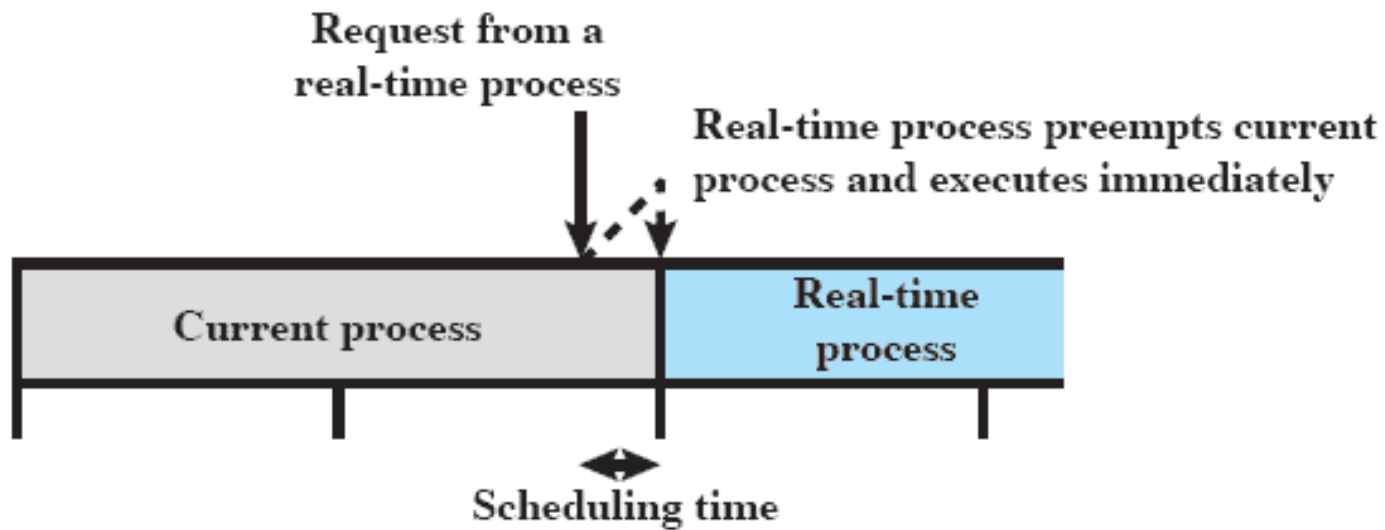
Scheduling of Real-Time Process



(c) Priority-Driven Preemptive Scheduler on Preemption Points

Still too slow for many real-time applications!

Scheduling of Real-Time Process



(d) Immediate Preemptive Scheduler

Real-Time Scheduling

- Static table-driven
 - Determines at run time when a task begins execution
- Static priority-driven preemptive
 - Traditional priority-driven scheduler is used
- Dynamic planning-based
 - Feasibility determined at run time
- Dynamic best effort
 - No feasibility analysis is performed

Deadline Scheduling

- Real-time applications are not concerned with speed but with completing tasks

Two Periodic Tasks

$T=(e,p)$

$A=(10,20)$

$B=(25,50)$

Table 10.2 Execution Profile of Two Periodic Tasks

| Process | Arrival Time | Execution Time | Ending Deadline |
|---------|--------------|----------------|-----------------|
| A(1) | 0 | 10 | 20 |
| A(2) | 20 | 10 | 40 |
| A(3) | 40 | 10 | 60 |
| A(4) | 60 | 10 | 80 |
| A(5) | 80 | 10 | 100 |
| • | • | • | • |
| • | • | • | • |
| • | • | • | • |
| B(1) | 0 | 25 | 50 |
| B(2) | 50 | 25 | 100 |
| • | • | • | • |
| • | • | • | • |
| • | • | • | • |

Scheduling with Completion DL

$T=(e,p)$
 $A=(10,20)$
 $B=(25,50)$

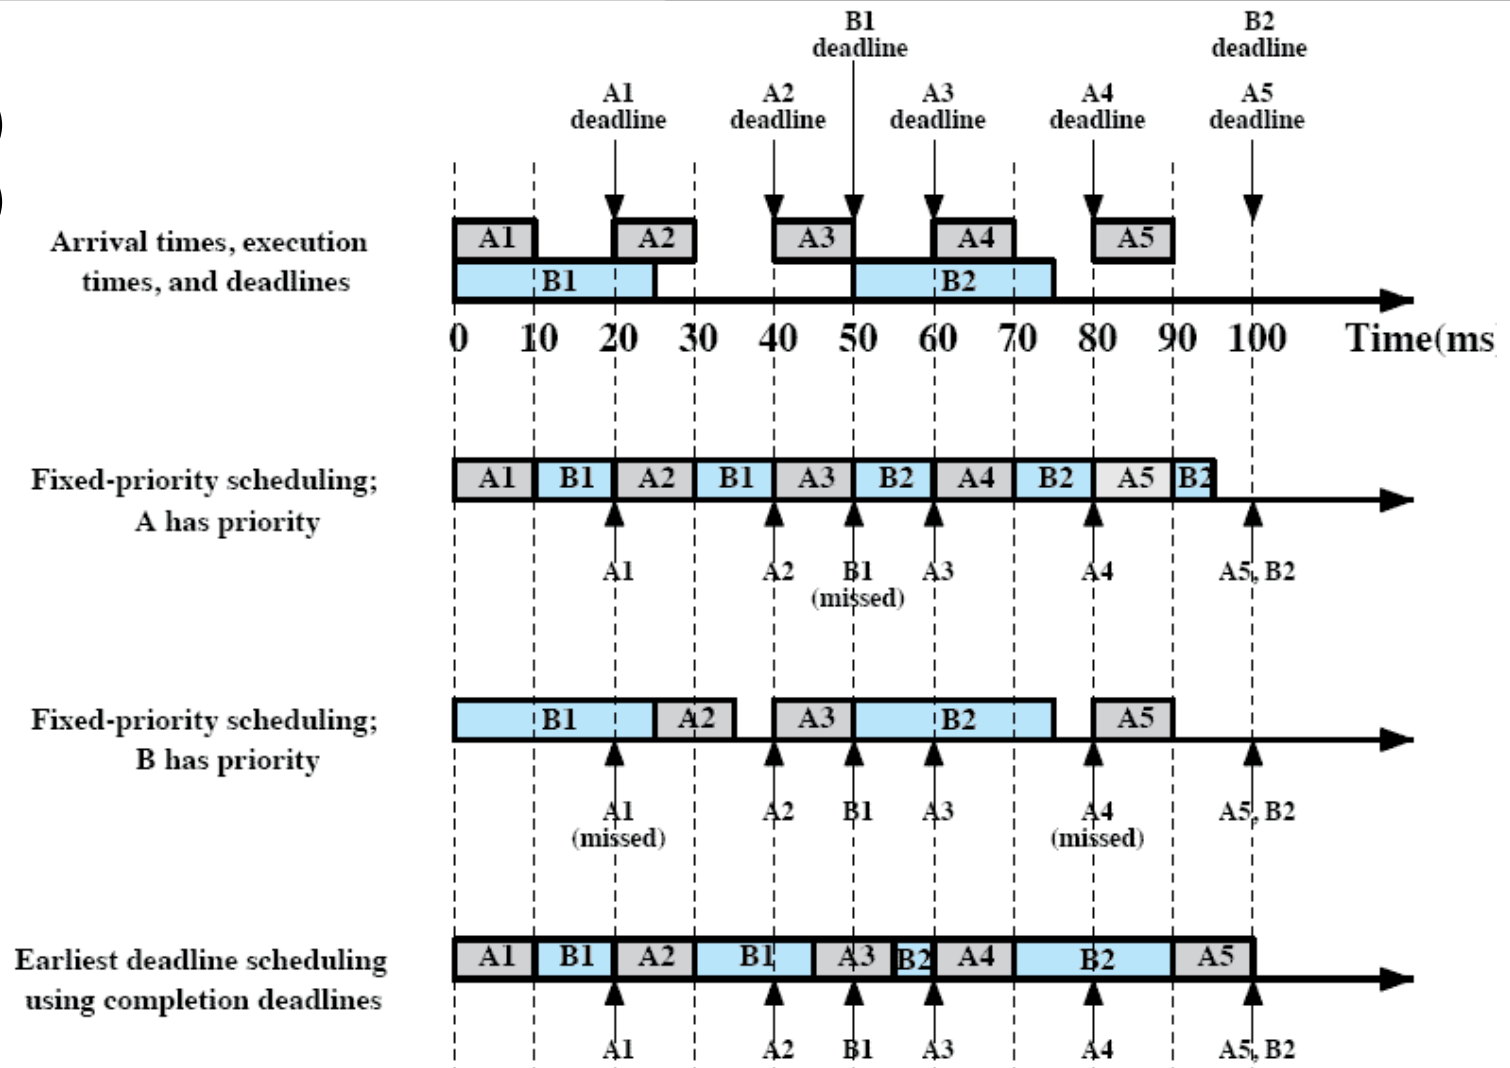


Figure 10.6 Scheduling of Periodic Real-time Tasks with Completion Deadlines (based on Table 10.2) 49

Execution Profile

Table 10.3 Execution Profile of Five Aperiodic Tasks

| Process | Arrival Time | Execution Time | Starting Deadline |
|----------------|---------------------|-----------------------|--------------------------|
| A | 10 | 20 | 110 |
| B | 20 | 20 | 20 |
| C | 40 | 20 | 50 |
| D | 50 | 20 | 90 |
| E | 60 | 20 | 70 |

Scheduling with Starting DL

No preemption

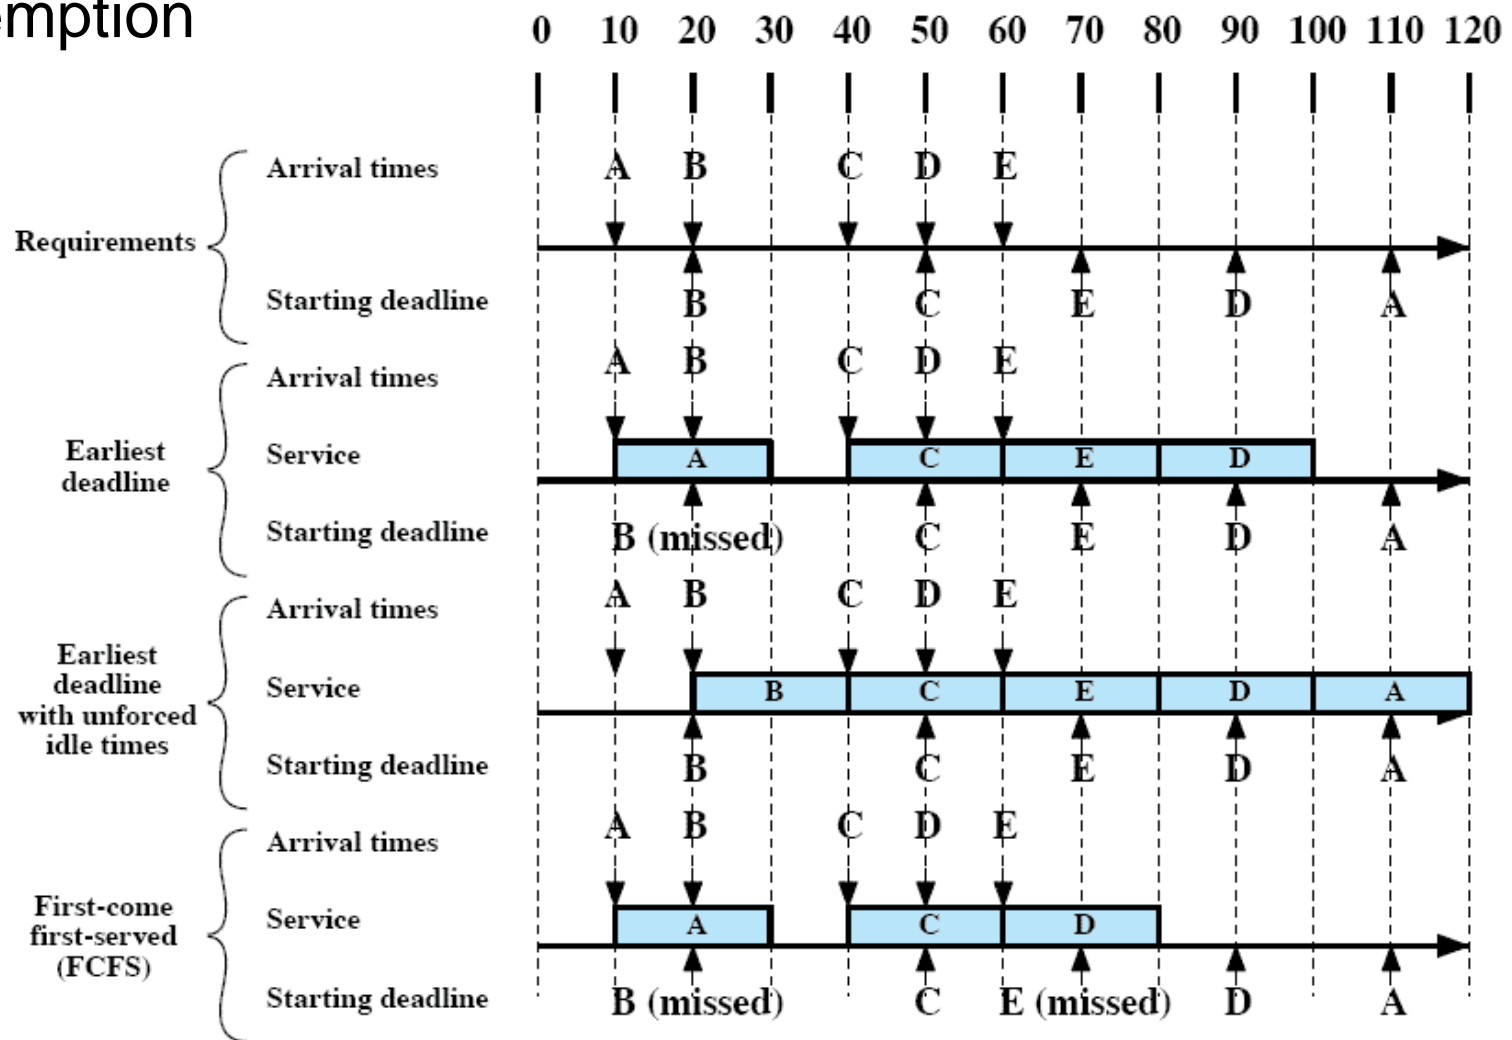


Figure 10.7 Scheduling of Aperiodic Real-time Tasks with Starting Deadlines

Rate Monotonic Scheduling

- Assigns priorities to tasks on the basis of their periods
- Highest-priority task is the one with the shortest period
- Lower bound on schedulable utilization = 0.693

$$\sum_{i=1}^n \frac{C_i}{T_i} \leq n(2^{\frac{1}{n}} - 1)$$

Task Set




Figure 10.8 A Task Set with RMS [WARR91]

Earliest Deadline First

- Always execute the task with the earliest deadline.
- Optimal (can schedule a CPU utilization of 1), if the system supports preemption

EDF vs. RM

... Or why does industry prefer RM?

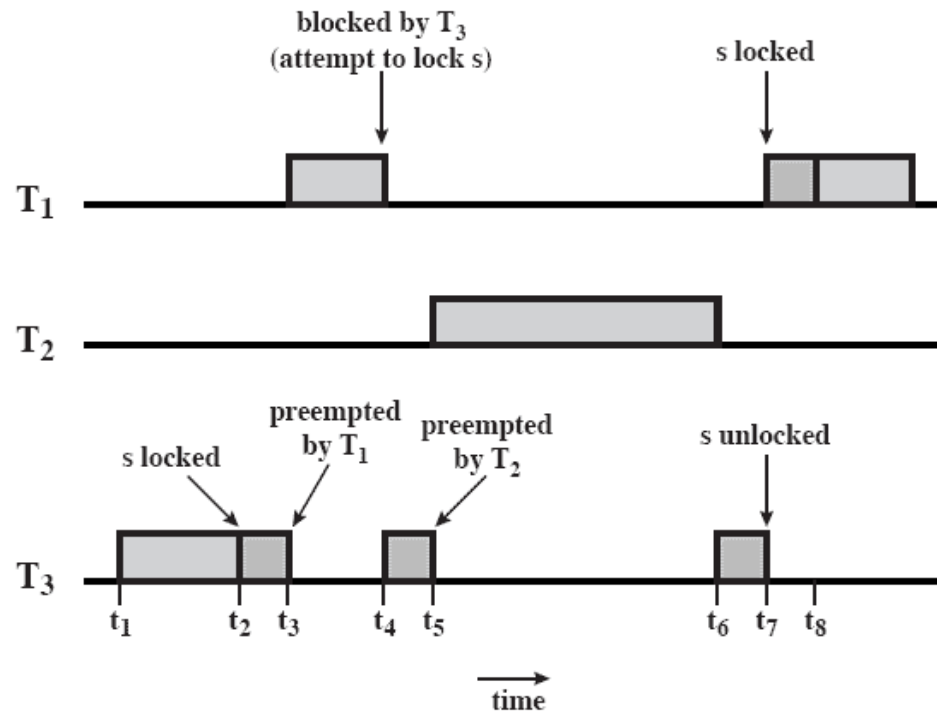
- 30% gain doesn't make much a difference
 - Processor speeds double
 - Good designs have a safety margin anyways
- Only parts are time critical
 - Use the ~30% for soft real-time tasks
- RM is more predictable in overload situations.
- Often you don't have deadlines.

Priority Inversion

- Can occur in any priority-based preemptive scheduling scheme
- Occurs when circumstances within the system force a higher priority task to wait for a lower priority task

Unbounded Priority Inversion

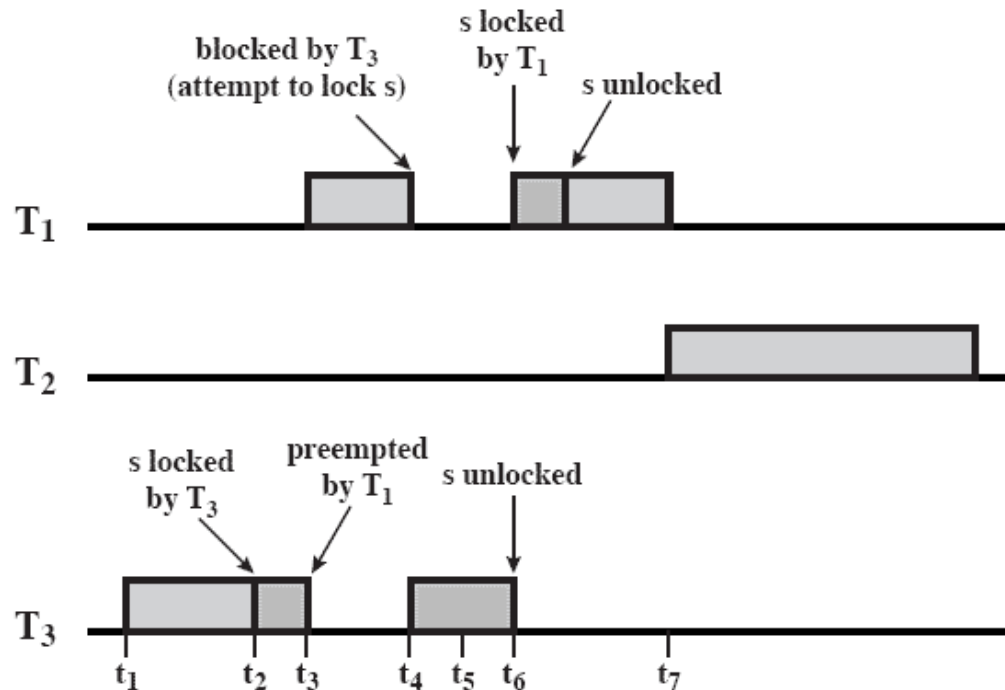
- Duration of a priority inversion depends on unpredictable actions of other unrelated tasks



(a) Unbounded priority inversion

Priority Inheritance

- Lower-priority task inherits the priority of any higher priority task pending on a resource they share



(b) Use of priority inheritance