

# **Chapter 11**

# **I/O Management and Disk Scheduling**

(based on original slides by Pearson)

# Categories of I/O Devices

---

- Human readable
  - Used to communicate with the user
  - Printers and terminals
- Machine readable
  - Used to communicate with electronic equipment
  - Disk drives, USB keys, Sensors, Controllers, Actuators
- Communication
  - Used to communicate with remote devices
  - Digital line drivers, modems

# Differences in I/O Devices

---

- Data rate
  - May be differences of several orders of magnitude between the data transfer rates

# I/O Device Data Rates

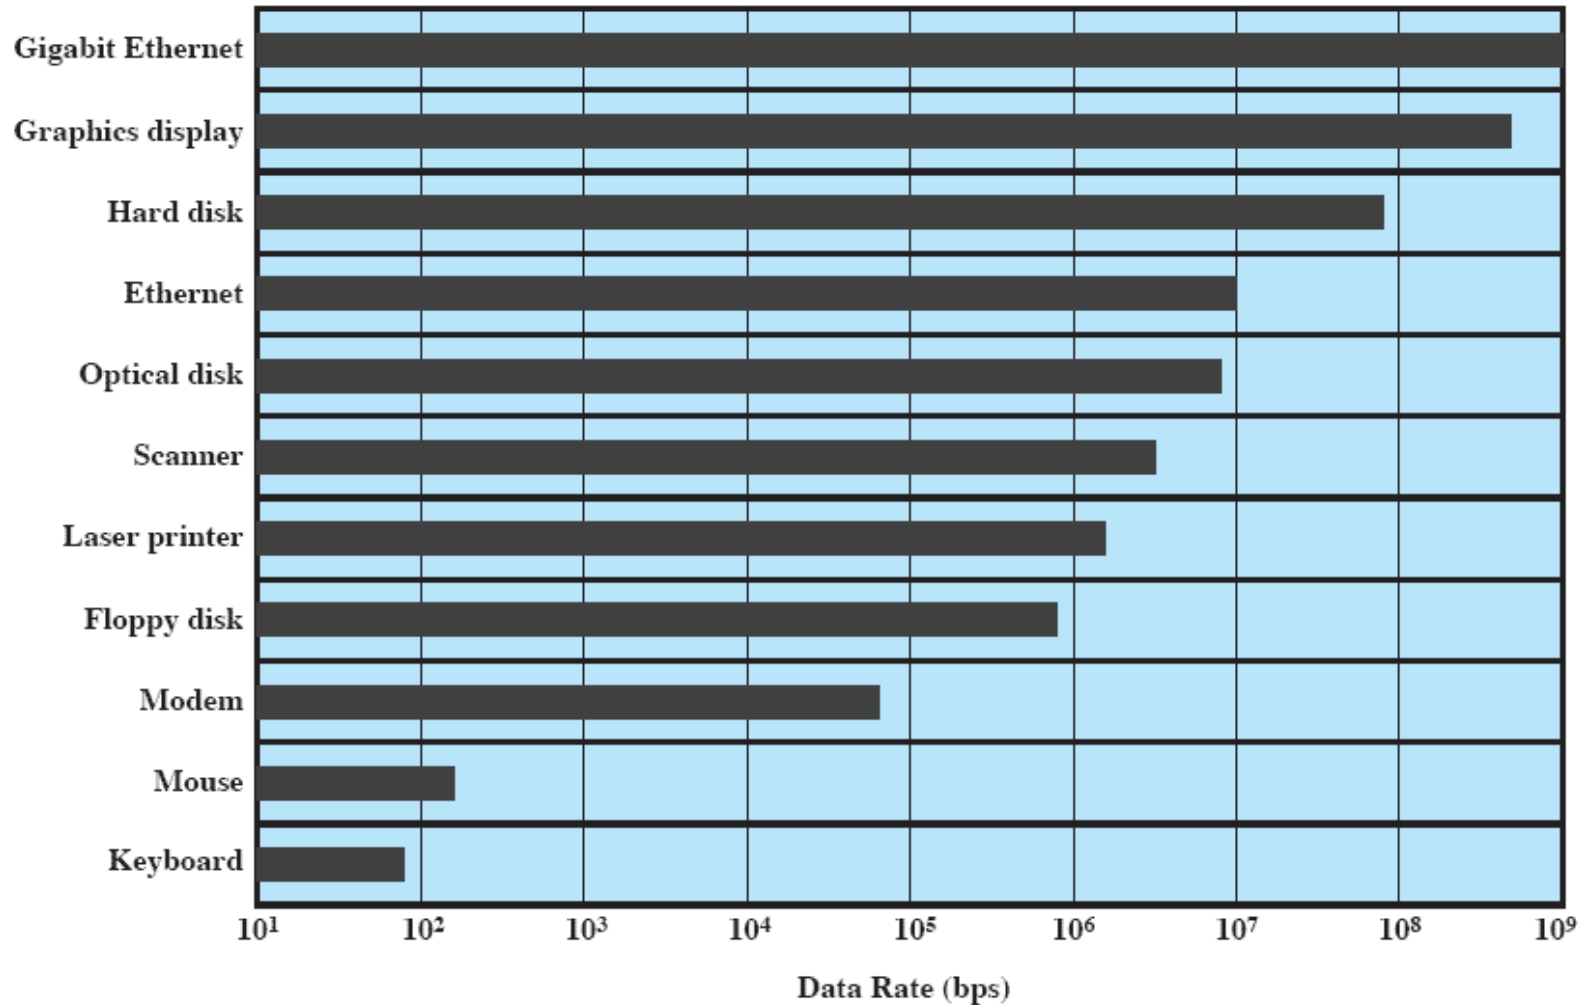


Figure 11.1 Typical I/O Device Data Rates

# Differences in I/O Devices

---

- Application
  - *The use of the device affects the management software*
  - Example:
    - Disk used to store files requires file management software
    - Disk used to store virtual memory pages needs special hardware and software to support it
    - Terminal used by system administrator may have a higher priority than the regular user's one

# Differences in I/O Devices

---

- Complexity of control
  - Simple one line on/off
  - Time dependent interaction: duty cycle of stepper motors
  - Complex, protocol-based interaction: most communication devices
- Unit of transfer
  - Data may be transferred as a **stream of bytes** for a terminal or in **larger blocks** for a disk

# Differences in I/O Devices

---

- Data representation
  - Encoding schemes, parity
- Error conditions
  - Devices respond to errors differently
    - Motor loses it's torque
    - CAN does simply not transmit

# Performing I/O

---

- Programmed I/O
  - Process is busy-waiting for the operation to complete
- Interrupt-driven I/O
  - I/O command is issued
  - Processor continues executing instructions
- Direct Memory Access (DMA)
  - DMA module controls exchange of data between main memory and the I/O device
  - Processor interrupted only after entire block has been transferred



# Relationship Among Techniques

---

**Table 11.1 I/O Techniques**

	<b>No Interrupts</b>	<b>Use of Interrupts</b>
<b>I/O-to-memory transfer through processor</b>	Programmed I/O	Interrupt-driven I/O
<b>Direct I/O-to-memory transfer</b>		Direct memory access (DMA)

# Evolution of the I/O Function

---

- Processor directly controls a peripheral device
- Controller or I/O module is added
  - Processor uses programmed I/O without interrupts
  - Processor does not need to handle details of external devices

# Evolution of the I/O Function

---

- Controller or I/O module with interrupts
  - Processor does not spend time waiting for an I/O operation to be performed
- Direct Memory Access
  - Blocks of data are moved into memory without involving the processor
  - Processor involved at beginning and end only

# Evolution of the I/O Function

---

- I/O module is a separate processor
- I/O processor
  - I/O module has its own local memory
  - It is a computer in its own right

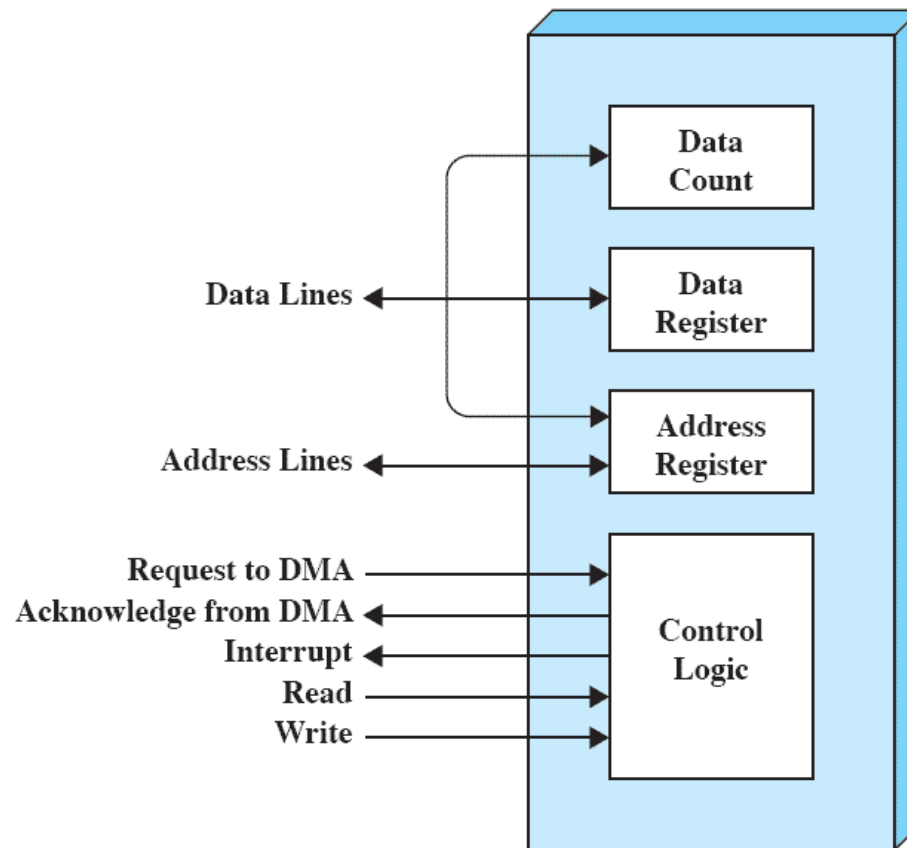
# Direct Memory Address

---

- Processor delegates I/O operation to the DMA module
- DMA module transfers data directly to or from memory
- When transfer is complete, DMA module sends an interrupt signal to the processor

# DMA

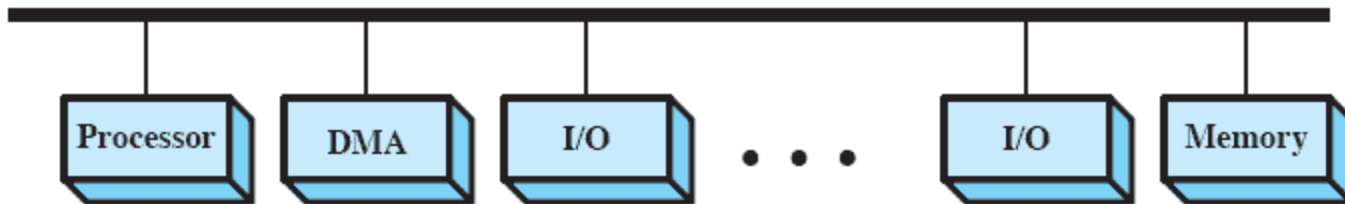
---



**Figure 11.2** Typical DMA Block Diagram

# DMA Configurations

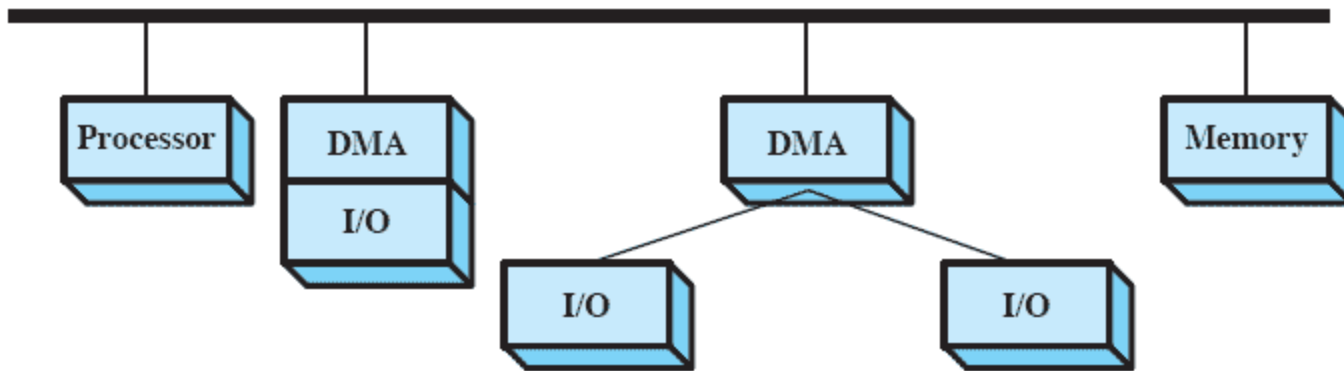
---



(a) Single-bus, detached DMA

# DMA Configurations

---

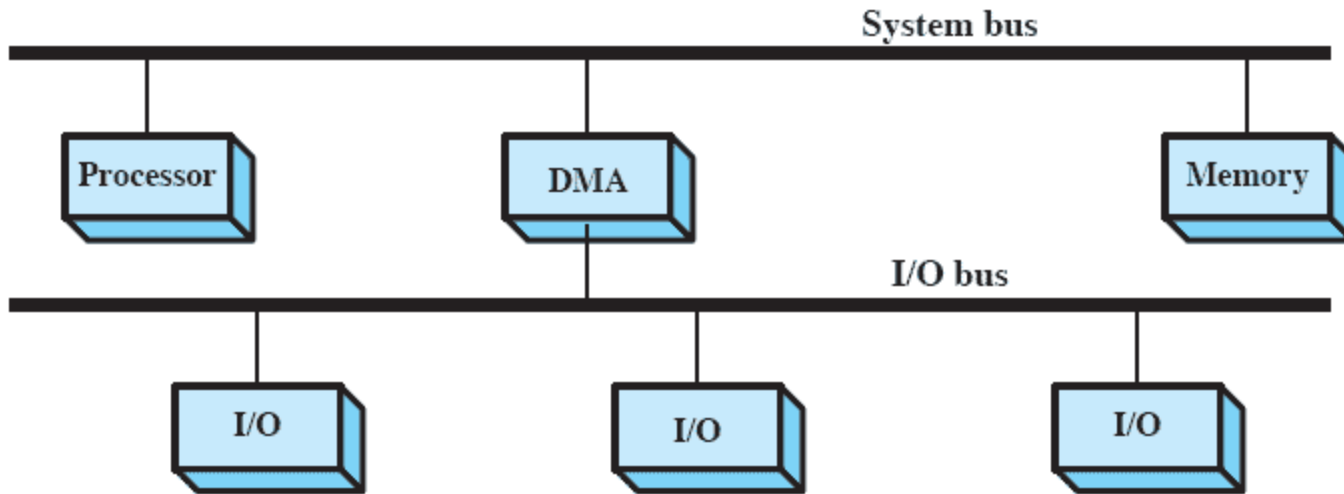


(b) Single-bus, Integrated DMA-I/O



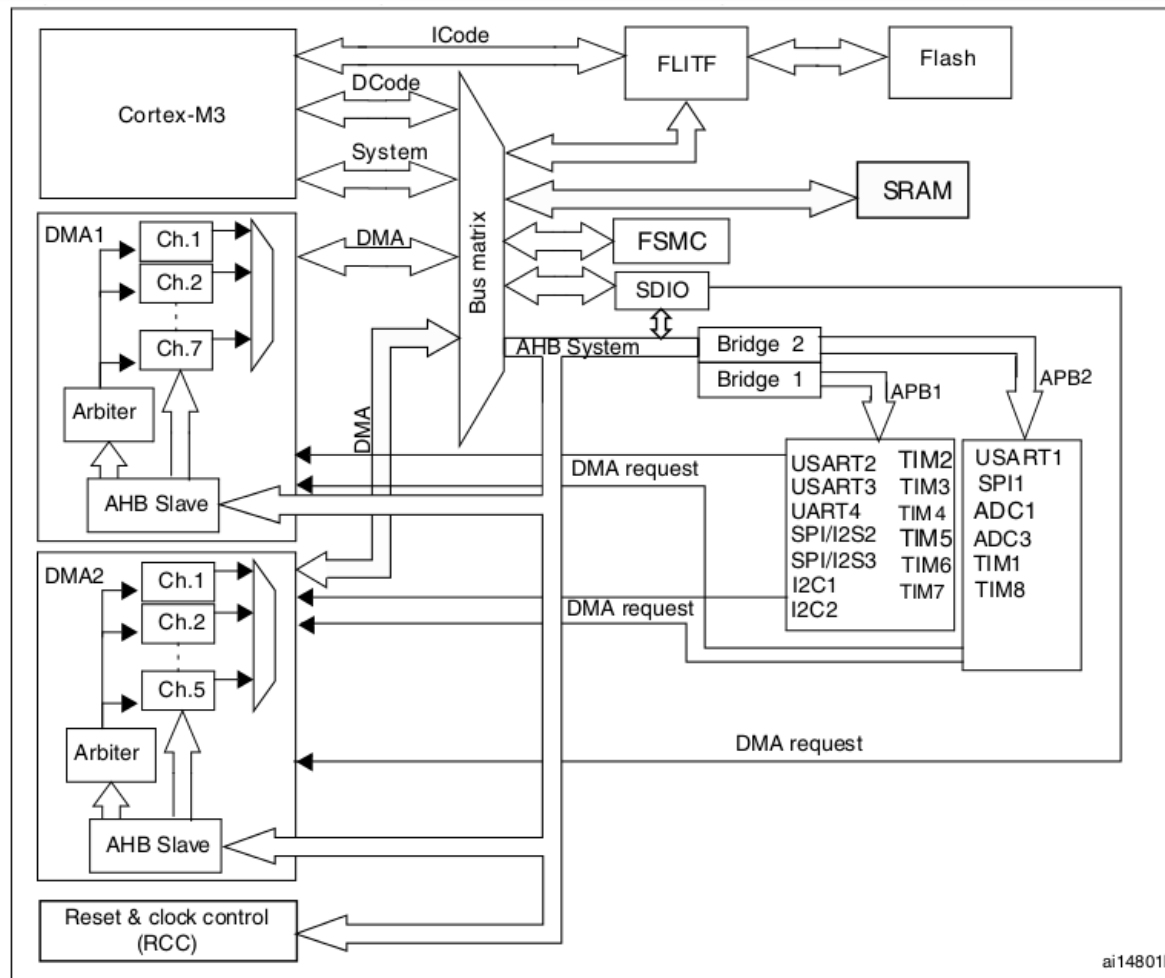
# DMA Configurations

---



(c) I/O bus

# DMA Real World



STM32F1xx MCU ~0.30\$/unit

# OS Design Issues wrt I/O

---

- Efficiency
  - Most I/O devices are extremely slow compared to main memory
  - Use of multiprogramming allows interleaving of I/O and processing
  - Even in the future: I/O will not keep up with processor throughputs
  - Swapping is used to bring in additional Ready processes; swapping is an I/O operation

# OS Design Issues wrt I/O

---

- Generality
  - Desirable to handle all I/O devices in a **uniform manner**
  - **Hide most of the details** of device I/O in lower-level routines

# I/O Buffering

---

- Reasons for buffering
  - Processes must wait for I/O to complete before proceeding
  - Certain pages must remain in main memory during I/O
- Risk of deadlock:
  - Process calls an I/O routine and blocks
  - Process is swapped out
  - Process waits for I/O completion, but I/O waits for the process to be swapped in to access the memory

# I/O Buffering

---

- Block-oriented
  - Information is stored in fixed sized blocks
  - Transfers are made a block at a time
  - Used for disks and USB keys
- Stream-oriented
  - Transfer information as a stream of bytes
  - Used for terminals, printers, communication ports, mouse and other pointing devices, and most other devices that are not secondary storage
- The kernel (I/O routines) handle buffering

# Single Buffer

---

- Operating system assigns a buffer in main memory for an I/O request
- Block-oriented
  - Input transfers made to buffer
  - Block moved to user space when needed
  - Another block is moved into the buffer
    - Read ahead

# Single Buffer

---

- Block-oriented
  - User process can process one block of data while next block is read in
  - Swapping can occur since input is taking place in system memory, not user memory
  - Operating system keeps track of assignment of system buffers to user processes



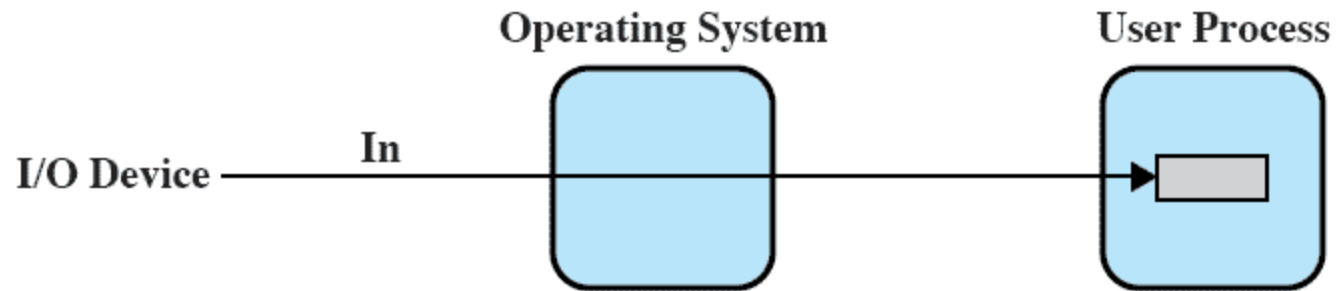
# Single Buffer

---

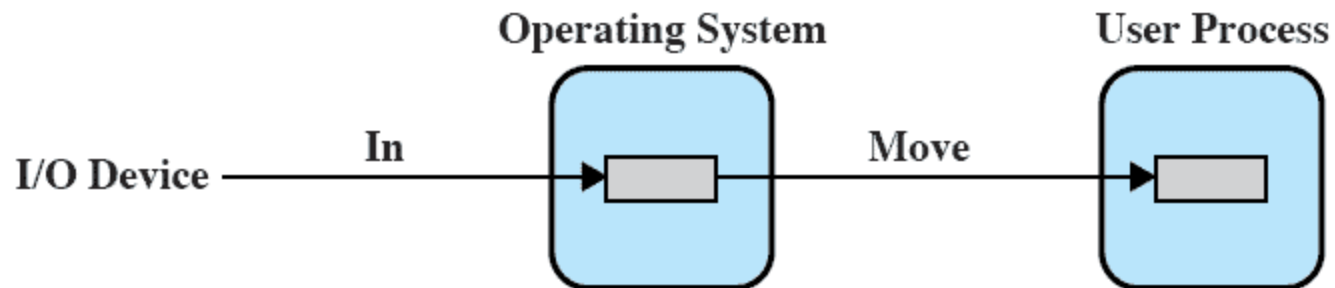
- Stream-oriented
  - Use a line at a time
  - User input from a terminal is one line at a time with carriage return signaling the end of the line
  - Output to the terminal is one line at a time
  - Flush() call in C

# Single Buffer

---



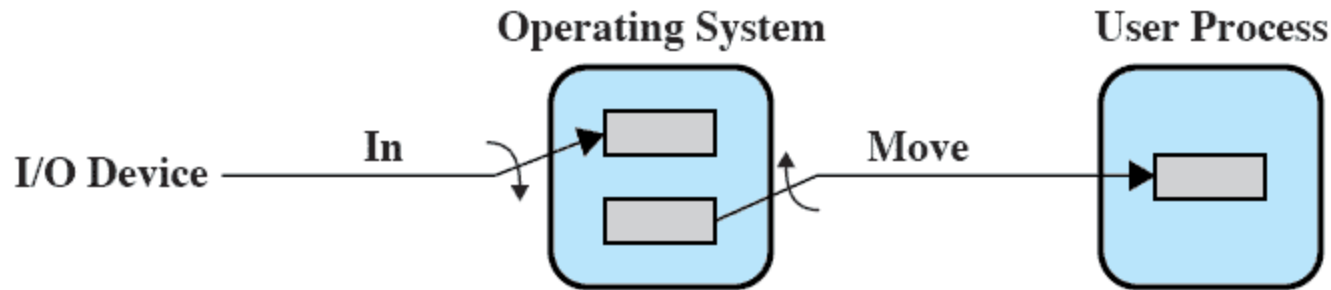
(a) No buffering



(b) Single buffering

# Double Buffer

- Use two system buffers instead of one
- A process can transfer data to or from one buffer while the operating system empties or fills the other buffer

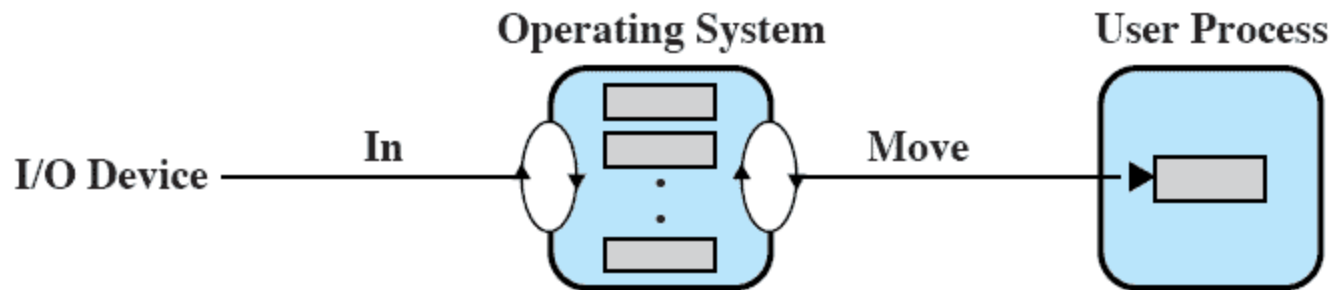


(c) Double buffering

- Process does not have to wait for I/O

# Circular Buffer

- More than two buffers are used
- Each individual buffer is one unit in a circular buffer



- Good for bursty I/O

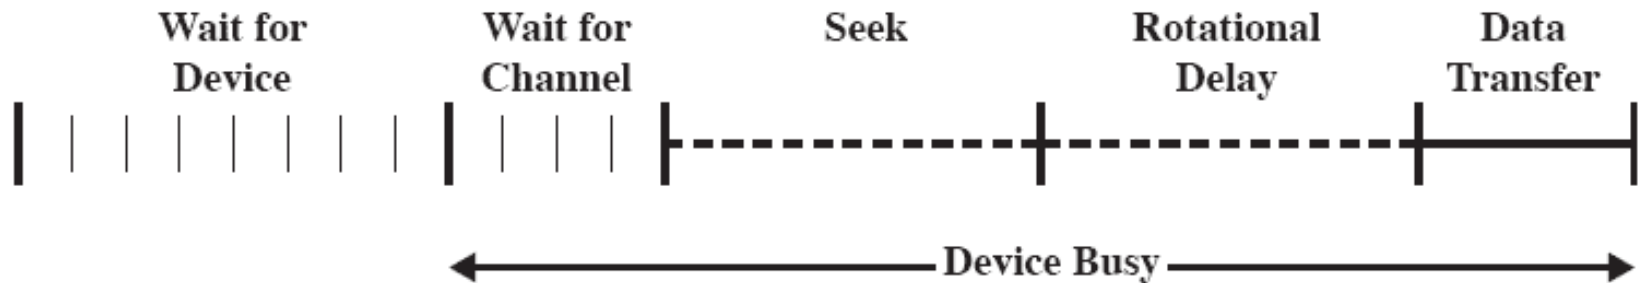
# Disk Performance Parameters

---

- To read or write, the disk head must be positioned at the desired track and at the beginning of the desired sector
- Seek time
  - Time it takes to position the head at the desired track
- Rotational delay or rotational latency
  - Time its takes for the beginning of the sector to reach the head

# Timing of Disk I/O Transfer

---



**Figure 11.6** Timing of a Disk I/O Transfer

# Disk Performance Parameters

---

- Access time
  - Sum of seek time and rotational delay
  - The time it takes to get in position to read or write
- There can be a huge difference between sequential read access and random read access.