


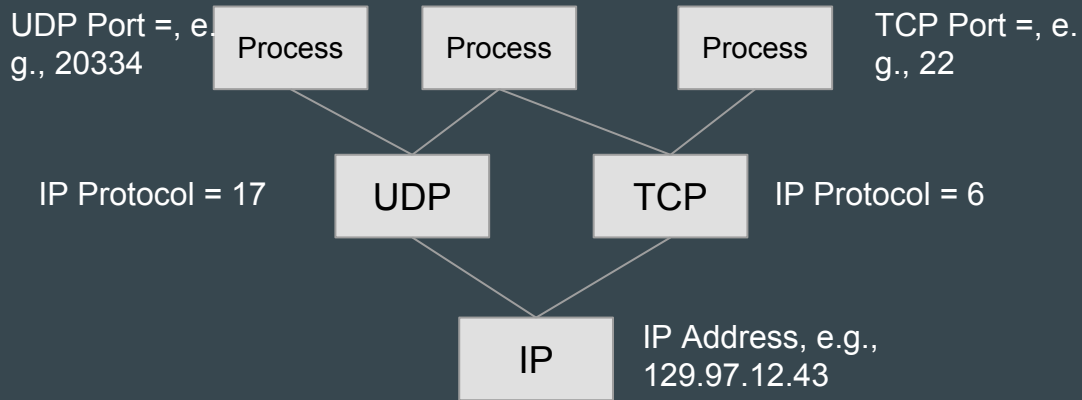
Focus on UDP & TCP over IP

- UDP & TCP are transport-layer protocols
- Over IP, which is a network-layer protocol




- UDP & TCP multiplexed over IP.
- Multiple processes multiplexed over each of UDP & TCP.

Multiplexing



The “5-tuple”

- On an (the) Internet, a connection or association is identified uniquely by the 5-tuple:
 - $\langle \text{source-ip-address, source-port, destination-ip-address, destination-port, protocol} \rangle$
 - E.g., $\langle 129.197.2.13, 20334, 216.58.199.14, 80, 6 \rangle$
 - E.g., $\langle 129.197.2.13, 50000, 216.58.199.14, 80, 6 \rangle$



The software view

- OS/library for UDP/TCP
- Applications can be written on top of UDP/TCP
- Client (initiator) - server (responder) paradigm.



The socket API

- POSIX standard API for UDP/TCP applications
- (Does not necessarily mean that it's a great API.)
- Essential calls:

- *socket()*
- *listen()*
- *accept()*
- *send()*
- *sendto()*

- *bind()*
- *connect()*
- *receive()*
- *sendto()*
- *recvfrom()*



UDP, TCP Applications →


Overlay Networking - p2p over TCP,UDP/IP


- Want a p2p network that maintains content.
- How to distribute content amongst peers?
- A client may contact some peer and want to lookup content.



The Chord DHT

- m -bit *key* unique for every peer and piece of content.
- Define: $\text{succ}(k)$ for any key k is peer that exists with smallest $\text{id} \geq k$.
- Content with key k hosted by $\text{succ}(k)$.






lookup(k)

- We may want to *lookup(k)* at any peer.
- Query is routed to *succ(k)*. How?



The Chord Approach

- Maintain a “finger table” (routing table) at Peer p , $FT_p[]$.
- m entries
- $FT_p[i] = succ(p + 2^{i-1})$, for all $i \in [1, m]$



Artwork credit

- ... (to be completed)