

# Cryptography

- What is cryptography?
- Related fields:
  - Cryptography (“**secret writing**”): Making secret messages
    - Turning **plaintext** (an ordinary readable message) into **ciphertext** (secret messages that are “hard” to read)
  - Cryptanalysis: Breaking secret messages
    - Recovering the plaintext from the ciphertext
- Cryptology is the science that studies these both
- The point of cryptography is to send secure messages over an insecure medium (like the Internet)

Comes from greek for secret writing. Cryptography is just making a plaintext message secret. Cryptography is a part of Cryptology which also includes cryptanalysis which is when we try to do the reverse.

# Dramatis Personae

- When talking about cryptography, we often use a standard cast of characters
- Alice, Bob, Carol, Dave
  - People (usually honest) who wish to communicate
- Eve
  - A passive eavesdropper, who can listen to any transmitted messages
- Mallory
  - An active Man-In-The-Middle, who can listen to, **and modify, insert, or delete**, transmitted messages
- Trent
  - A Trusted Third Party

We standardize our fictional characters.

- Alice, Bob, Carol, Dave - good honest hacker
- Eve - a passive eavesdropper that can't do anything
- Mallory - man in the middle
- Trent - trusted third party

# Building blocks

- Cryptography contains three major types of components
  - Confidentiality components
    - Preventing Eve from **reading** Alice's messages
  - Integrity components
    - Preventing Mallory from **modifying** Alice's messages without being detected
  - Authenticity components
    - Preventing Mallory from **impersonating** Alice

Confidentiality means we want to prevent Eve from reading Alice's messages. We can do this using a cryptosystem

Integrity says we want to prevent Mallory from modifying Alice's messages. We can do this through message authentication (MACs), signature schemes, or cryptographic hash functions (not very secure since hackers can do this easily).

Privacy says we want to prevent Mallory from impersonating Alice. We can do this by passwords or biometrics and such, or through challenge response protocols.

A secret-key system or a public-key system. With a public key system we have two keys. A public key that anyone can use to encrypt something, and a secret private key that can be used to decrypt stuff. A secret key system can be block (requiring data to be a fixed width) or stream.

# Kerckhoffs' Principle (19th c.)

- The security of a cryptosystem should not rely on a secret that's hard (or expensive) to change
- So don't have secret encryption methods
  - Then what do we do?
  - Have a large class of encryption methods, instead
    - Hopefully, they're all equally strong
  - Make the class **public** information
  - Use a secret **key** to specify which one you're using
  - It's easy to change the key; it's usually just a smallish number

We should only ever rely on a simple secret key that is easy to change. We could have a bunch of different encryption functions. This is not always practical so we encrypt with a long key that varies each time. A system is only as secure as the possible number of keys (we want the widest variety of keys to prevent brute forcing). Even then it is always possible to brute force no matter how long your key is.



# Kerckhoffs' Principle (19th c.)

- This has a number of implications:
  - The system is at **most** as secure as the number of keys
  - Eve can just try them all, until she finds the right one
  - A **strong cryptosystem** is one where that's the best Eve can do
    - With weaker systems, there are shortcuts to finding the key
  - Example: newspaper cryptogram has 403,291,461,126,605,635,584,000,000 possible keys
  - But you don't try them all; it's way easier than that!

A **strong crypto system** is one where iterating through every possible key is the best possible attack.


# Strong cryptosystems

- What information do we assume the attacker (Eve) has when she's trying to break our system?
- She may:
  - Know the **algorithm** (the public class of encryption methods)
  - Know some **part of the plaintext**
  - Know a number (maybe a large number) of corresponding **plaintext/ciphertext pairs**
  - Have access to an encryption and/or decryption **oracle**
- And we still want to prevent Eve from learning the key!

Frequently the attacker can get a small bit of plaintext from cypher text. For instance HTTP headers tend to be fairly standard so intercepting an encrypted request allows you to figure out a little bit. Breaking the enigma machine was done by getting a bunch of plaintext cyphertext pairs and analyzing them for a bit before breaking the code. To get these plain/cypher pairs they listened to this one base in russia that would send the same message every day (“weather report: all clear”).

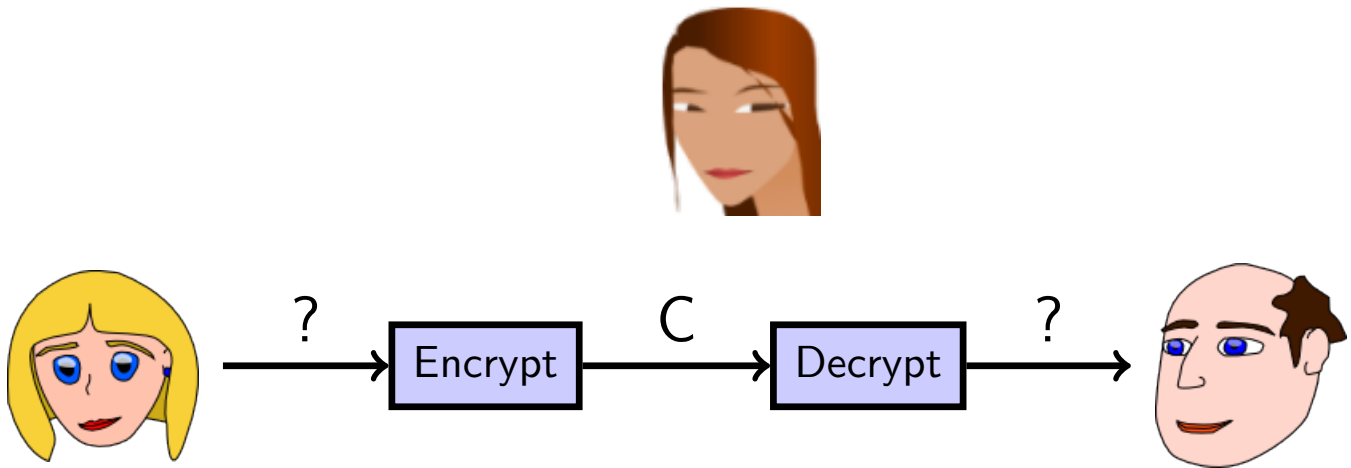
# Secret-key encryption

- Secret-key encryption is the simplest form of cryptography
- Also called symmetric encryption
- Used for thousands of years
- The key Alice uses to encrypt the message is the same as the key Bob uses to decrypt it



# Secret-key encryption

- Eve, not knowing the key, should not be able to recover the plaintext



# Perfect secret-key encryption

- Is it possible to make a completely unbreakable cryptosystem?
- Yes: the **One-Time Pad**
- It's also very simple:
  - The key is a truly random bitstring of the same length as the message
  - The “Encrypt” and “Decrypt” functions are each just XOR

# One-time pad

- Q: Why does “try every key” not work here?
- It's very hard to use correctly
  - The key must be **truly random**, not pseudorandom
  - The key must **never be used more than once!**
    - A “two-time pad” is **insecure!**
- Q: How do you share that much secret key?
- Used in the Washington / Moscow hotline for many years



# Computational security

- In contrast to OTP's "perfect" or "information-theoretic" security, most cryptosystems have "computational" security
  - This means that it's certain they can be broken, given enough work by Eve
- How much is "enough"?
- At **worst**, Eve tries every key
  - How long that takes depends on how long the keys are
  - But it only takes this long if there are no "shortcuts"!

## Some data points

- One computer can try about 17 million keys per second
- A medium-sized corporate or research lab may have 100 computers
- The BOINC project has 13 million computers



Berkeley Open Infrastructure  
for Network Computing

- Remember that most computers are idle most of the time (they're waiting for you to type something); getting them to crack keys in their spare time doesn't actually cost anything extra

## 40-bit crypto

- This was the US legal export limit for a long time
- $2^{40} = 1,099,511,627,776$  possible keys
- One computer: 18 hours
- One lab: 11 minutes
- BOINC: 5 ms

A guy came up with pgp which allowed any key length. He ended up getting sued by the US government because encryption was classified as munition.

## 56-bit crypto

- This was the US government standard (DES) for a long time
- $2^{56} = 72,057,594,037,927,936$  possible keys
- One computer: 134 years
- One lab: 16 months
- BOINC: 5 minutes

A comity came up with the stupid number of 56 because the government wanted 48 so that they could break it with super computers but people wanted 64 so it would be more secure, so they split the difference.

# Cracking DES



“DES cracker” machine of Electronic Frontier Foundation

People wanted to show that DES was not all that secure so they held a contest for people to try to break it. Eventually the Electronic Frontier Foundation did it using a special machine. It took about 2 days to crack a key. Researchers eventually made a system called copacabana that used fpgas (120, cost 10k) that could crack DES in about a week.



# 128-bit crypto

- This is the modern standard
- $2^{128} = 340,282,366,920,938,463,463,374,607,431,768,211,456$  possible keys
- One computer: 635 thousand million million million years
- One lab: 6 thousand million million million years
- BOINC: 49 thousand million million years

Nowadays we use 128 bits. This is so huge that you cannot brute force attack it. So if the encryption is secure it is “unhackable”. That being said, computers are continually getting faster by moore’s law. So in 132 years you can break it in a day.

## Well, we cheated a bit

- This isn't really true, since computers get faster over time
  - A better strategy for breaking 128-bit crypto is just to wait until computers get  $2^{88}$  times faster, then break it on one computer in 18 hours.
  - How long do we wait? Moore's law says 132 years.
  - If we believe Moore's law will keep on working, we'll be able to break 128-bit crypto in 132 years (and 18 hours) :-)
    - Q: Do we believe this?

## An even better strategy

- Don't break the crypto at all!
- There are always weaker parts of the system to attack
  - Remember the Principle of Easiest Penetration
- The point of cryptography is to make sure the information transfer is not the weakest link

## Rubber hose cryptanalysis

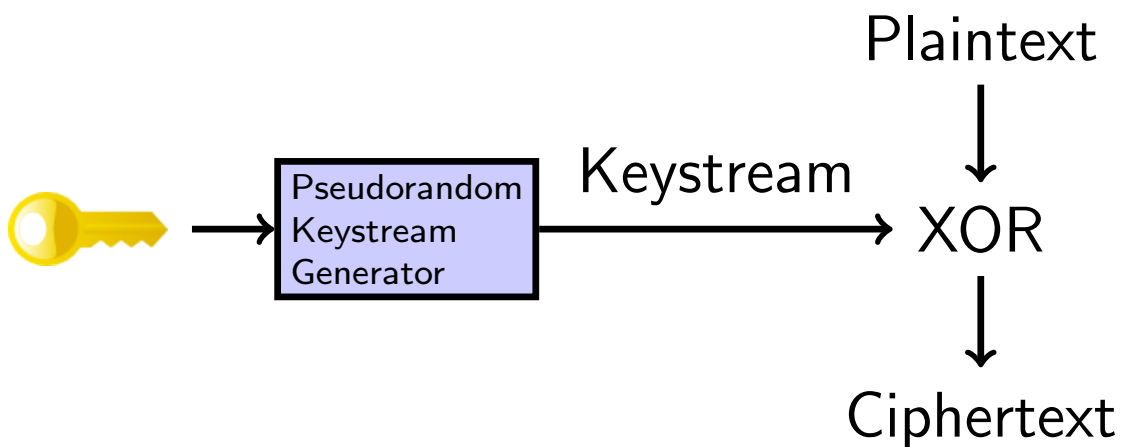


# Types of secret-key cryptosystems

- Secret-key cryptosystems come in two major classes
  - Stream ciphers
  - Block ciphers

# Stream ciphers

- A stream cipher is what you get if you take the One-Time Pad, but use a pseudorandom keystream instead of a truly random one



- **RC4** is the most commonly used stream cipher on the Internet today

RC4 is the most commonly used stream cipher, but it has a ton of problems (new hacks every year).



# Stream ciphers

- Stream ciphers can be very fast
  - This is useful if you need to send a **lot** of data securely
- But they can be tricky to use correctly!
  - What happens if you use the same key to encrypt two different messages?
  - How would you solve this problem without requiring a new shared secret key for each message? Where have we seen this technique before?
- WEP, PPTP are great examples of how **not** to use stream ciphers

Stream ciphers use the  $\otimes$  operator since computers are super fast at them. They work using salts, but you have to be very sure that the salts are unique. Take your plaintext,  $\otimes$  with key and add the salt. The encrypted text is sent along with the salt. Important to note, the salt is public knowledge.

# Block ciphers

- Note that stream ciphers operate on the message one bit at a time
- What happens in a stream cipher if you change just one bit of the plaintext?
- We can also use block ciphers
  - Block ciphers operate on the message one block at a time
  - Blocks are usually 64 or 128 bits long
- **AES** is the block cipher everyone should use today
  - Unless you have a really, really good reason

If we know the location in the string of the value you want to fuck with you can flip bits and such to change that number.


Most block ciphers (including AES) are based on substitution-permutation networks. They repeat their encoding steps a number of times around and around.

1.  $\otimes$  key bit state
2. apply s-boxes (substitution box) which just takes a box of data and map it to something else (usually every 4 bits)
3. permute bits

At some point it was revealed that by switching to using the NSA's sbox mapping DES became more resistant to attacks around sboxes. This showed that they were about 10 years ahead of academic researchers.

# Modes of operation

- Block ciphers work like this:




- But what happens when the plaintext is larger than one block?
  - The choice of what to do with multiple blocks is called the **mode of operation** of the block cipher

We have a block of plaintext, put it through cipher and get garbled mess. Problem arises when the block of plaintext is not the correct length (block must be multiple of some number).

# Modes of operation

- The simplest thing to do is just to encrypt each successive block separately.
  - This is called Electronic Code Book (**ECB**) mode
- But if there are repeated blocks in the plaintext, you'll see the same repeating patterns in the ciphertext:



Take your message and break it up into a bunch of fixed widths and just encrypt each one. This isn't that good because if any of those blocks are the same they will be encrypted the same. These patterns will appear in the encrypted text giving away more information than was intended. Think of a black and white image.

Never use ECB mode for encryption because it does this.



# Modes of operation

- There are much better modes of operation to choose from
  - Common ones include Cipher Block Chaining (**CBC**), Counter (**CTR**), and Galois Counter (**GCM**) modes
- Patterns in the plaintext are no longer exposed
- But you need an **IV** (Initial Value), which acts much like a salt



A solution for this is to make each chunk dependent on the other chunks. Take the cipher text from the previous block and and it with the key for the next block to use as the key for that block. The message becomes a block longer because you need some starting cipher text to include.

# Key exchange

- How do Alice and Bob share the secret key?
  - Meet in person; diplomatic courier
  - In general this is very hard
- Or, we invent new technology...

There is a ton of key information flying around so we need a good way to get it to the people who need them.

# Public-key cryptography


- Invented (in public) in the 1970's
- Also called asymmetric cryptography
  - Allows Alice to send a secret message to Bob **without** any prearranged shared secret!
  - In secret-key cryptography, the same (or a very similar) key encrypts the message and also decrypts it
  - In public-key cryptography, there's one key for encryption, and a **different** key for decryption!
- Some common examples:
  - RSA, ElGamal, ECC, NTRU, McEliece

# Public-key cryptography

- How does it work?
  - Bob gives everyone a copy of his public encryption key. Alice uses it to encrypt a message, and sends the encrypted message to Bob
  - Bob uses his private decryption key to decrypt the message.
    - Eve can't decrypt it; she only has the encryption key.
    - Neither can Alice!
- So with this, Alice just needs to know Bob's public key in order to send him secret messages
  - These public keys can be published in a directory somewhere

One key lets you encrypt and a different one lets you decrypt. An example is RSA. Public key = everyone knows it, secret key = no one knows it

# Public-key cryptography





This is hackable if Mallory publishes a key to Alice to replace Bob's key with her own. When Alice looks up a public key Mallory replaces the encryption key with her own. Then she knows the decryption key so she can just listen in on the conversation.

RSA was published in 1977 by Clifford Cox. It works by generating two large prime numbers,  $p$  and  $q$  that are inverses of each other. So  $n = pq$  and  $ab \equiv 1 \pmod{(p-1)(q-1)}$ . The public key is then  $(n, a)$  and the private key is  $(p, q, b)$ . Encryption is  $x^a \bmod n$  and decryption is  $x^b \bmod n$ . You want these numbers to be huge, but this makes doing the calculations very slow.

It is important that these are prime numbers else you could just factor the modulus.

# Public key sizes

- Recall that if there are no shortcuts, Eve would have to try  $2^{128}$  things in order to read a message encrypted with a 128-bit key
- Unfortunately, all of the public-key methods we know **do** have shortcuts
  - Eve could read a message encrypted with a 128-bit RSA key with just  $2^{33}$  work, which is **easy**!
  - If we want Eve to have to do  $2^{128}$  work, we need to use a much longer public key

Public key systems have to have very large keys because people know them.

# Public key sizes

Comparison of key sizes for roughly equal strength

<u>AES</u>	<u>RSA</u>	<u>ECC</u>
80	1024	160
116	2048	232
128	2600	256
160	4500	320
256	14000	512

To have equivalent security parameters the public key crypto will take 3 orders of magnitude more time to encryption than a symmetric crypto system.

# Hybrid cryptography

- In addition to having longer keys, public-key cryptography takes a long time to calculate (as compared to secret-key cryptography)
  - Using public-key to encrypt large messages would be too slow, so we take a hybrid approach:
    - Pick a random 128-bit key  $K$  for a secret-key cryptosystem
    - Encrypt the large message with the key  $K$  (e.g., using AES)
    - Encrypt the key  $K$  using a public-key cryptosystem
    - Send the encrypted message and the encrypted key to Bob
  - This hybrid approach is used for almost every cryptography application on the Internet today

To get the best of both worlds we combine symmetric and public encryption. Say we have some plain text  $P$  that is a very long message. We generate an AES key called the session key. We get cipher text  $C$  by encrypting  $P$  using the session key. We then encrypt the session key using Bob's public key using a public key system like RSA.

## Is that all there is?

- It seems we've got this "sending secret messages" thing down pat. What else is there to do?
  - Even if we're safe from Eve reading our messages, there's still the matter of Mallory
  - It turns out that even if our messages are encrypted, Mallory can sometimes modify them in transit!
  - Mallory won't necessarily know what the message says, but can still change it in an undetectable way
    - e.g. bit-flipping attack on stream ciphers
  - This is counterintuitive, and often forgotten
    - The textbook even gets this wrong!
- How do we make sure that Bob gets the same message Alice sent?



We still run into integrity problems. Just because it is secure does not mean that the message contents have not been altered in path.

# Integrity components

- How do we tell if a message has changed in transit?
- Simplest answer: use a checksum
  - For example, add up all the bytes of a message
  - The last digits of serial numbers (credit card, ISBN, etc.) are usually checksums
  - Alice computes the checksum of the message, and sticks it at the end before encrypting it to Bob. When Bob receives the message and checksum, he verifies that the checksum is correct

# This doesn't work!

- With most checksum methods, Mallory can easily change the message in such a way that the checksum stays the same
- We need a “cryptographic” checksum
- It should be hard for Mallory to find a second message with the same checksum as any given one

Checksums are usually very simple so they really only work for accidentally errors, most attackers will just compute the checksum and set it properly,

Hash functions are nice because they don't have any key at all.

Authentication codes (mentioned earlier as MACs) have a shared secret key.

Signature schemes have a private signing algorithm and a public verification algorithm.

# Cryptographic hash functions

- A **hash function**  $h$  takes an arbitrary length string  $x$  and computes a fixed length string  $y = h(x)$  called a **message digest**
  - Common examples: MD5, SHA-1, SHA-2, SHA-3 (AKA Keccak, from 2012 on)
- Hash functions should have three properties:
  - Preimage-resistance:
    - Given  $y$ , it's hard to find  $x$  such that  $h(x) = y$  (i.e., a “preimage” of  $x$ )
  - Second preimage-resistance:
    - Given  $x$ , it's hard to find  $x' \neq x$  such that  $h(x) = h(x')$  (i.e., a “second preimage” of  $h(x)$ )
  - Collision-resistance:
    - It's hard to find any two distinct values  $x, x'$  such that  $h(x) = h(x')$  (a “collision”)

MD5 is shit, do not use.

SHA is a set of standards published that hash algorithms should follow. Currently there have been no fully breaking attacks on SHA-1 just weakening ones, we want to all be moving to SHA-3.

**Preimage resistance** you should not be able to find any input that gives this output.

**Second preimage resistance** you should not be able to find another input that gives this output

**Collision resistance** you shouldn't be able to find an input or its output

We really want to make is so that hackers cannot find any collisions (when we accomplish this it is a strong hash function).

# What is “hard” ?

- For SHA-1, for example, it takes  $2^{160}$  work to find a preimage or second preimage, and  $2^{80}$  work to find a collision using a brute-force search
  - However, there are faster ways than brute force to find collisions in SHA-1 or MD5
- Collisions are always easier to find than preimages or second preimages due to the well-known birthday paradox

Preimage resistance has us think of a bunch of inputs and try to get the given output of  $y$  (do the same for second preimage). Collision resistance is much easier since you just want to see if any two inputs have matching outputs, so it takes a square root number of times that preimage checks need.

**Birthday Paradox** The odds of two people sharing a birthday in a room are significantly higher than the odds of someone in the room having the same birthday as me. This only takes roughly 23 people to get 50% probability of two sharing a birthday.

You can weaken these hashes if they preserve character frequencies.



# What is “hard” ?

- For SHA-1, for example, it takes  $2^{160}$  work to find a preimage or second preimage, and  $2^{80}$  work to find a collision using a brute-force search
  - However, there are faster ways than brute force to find collisions in SHA-1 or MD5
- Collisions are always easier to find than preimages or second preimages due to the well-known birthday paradox

Currently this does not preserve integrity because Mallory knows the hash function so she can just rewrite the message, hash it, and swap that in.

# Cryptographic hash functions

- You can't just send an unencrypted message and its hash to get integrity assurance
  - Even if you don't care about confidentiality!
- Mallory can change the message and just compute the new message digest herself

# Cryptographic hash functions

- Hash functions provide integrity guarantees only when there is a secure way of sending and/or storing the message digest
  - For example, Bob can publish a hash of his public key (i.e., a message digest) on his business card
  - Putting the whole key on there would be too big
  - But Alice can download Bob's key from the Internet, hash it herself, and verify that the result matches the message digest on Bob's card
- What if there's no external channel to be had?
  - For example, you're using the Internet to communicate

So hash functions are only good for preserving integrity when they cannot be messed with.

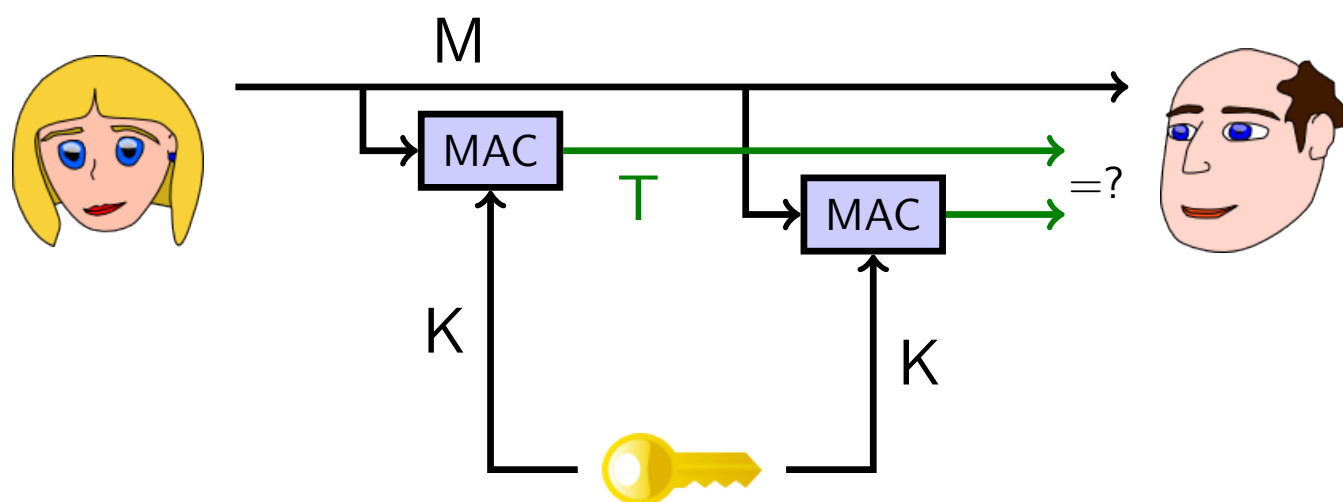
# Message authentication codes

- We do the same trick as for encryption: have a large class of hash functions, and use a shared secret key to pick the “correct” one
- Only those who know the secret key can generate, or even check, the computed hash value (sometimes called a **tag**)
- These “keyed hash functions” are usually called **Message Authentication Codes**, or **MACs**
- Common examples:
  - SHA-1-HMAC, SHA-256-HMAC, CBC-MAC

When we encrypt the message we take the last block of cipher text and that becomes the MAC

**HMAC** returns the hash of (key  $\otimes$  opad append hash of (key  $\otimes$  ipad append message)). opad and ipad are fixed public settings.

# Message authentication codes





Alice wants to get message to Bob so she sends it using the shared secret key  $k$ . Alice also sends him the MAC of the message, from this Bob calculates the MAC of the message that Alice sent him and checks that this matches the MAC she sent him. Because the MAC calculation requires knowledge of the secret key you can use this to ensure integrity.

# Combining ciphers and MACs

- In practice we often need both confidentiality and message integrity
- There are multiple strategies to combine a cipher and a MAC when processing a message
  - Encrypt-then-MAC, MAC-then-Encrypt, Encrypt-and-MAC
- Encrypt-then-MAC is the recommended strategy
- Ideally your crypto library already provides an **authenticated encryption mode** that securely combines the two operations so you don't have to worry about getting it right
  - E.g., GCM, CCM (used in WPA2, see later), or OCB mode

**Encrypt then MAC** You can encrypt the message first and then MAC it, return `encrypt(message) append MAC(encrypt(message))`

**MAC then encrypt** `encrypt(message append MAC(message))`

**encryot and MAC** `encrypt(message) append MAC(message)`

Encrypt then MAC is the best of the options.

# Repudiation

- Suppose Alice and Bob share a MAC key  $K$ , and Bob receives a message  $M$  along with a valid tag  $T$  that was computed using the key  $K$ 
  - Then Bob can be assured that Alice is the one who sent the message  $M$ , and that it hasn't been modified since she sent it!
  - This is like a “signature” on the message
  - But it's not quite the same!
  - Bob can't show  $M$  and the tag  $T$  to Carol to prove Alice sent the message  $M$

# Repudiation

- Alice can just claim that Bob made up the message  $M$ , and calculated the tag  $T$  himself
- This is called **repudiation**; and we sometimes want to avoid it
- Some interactions should be repudiable
  - Private conversations
- Some interactions should be non-repudiable
  - Electronic commerce

# Digital signatures

- For non-repudiation, what we want is a true **digital signature**, with the following properties:
- If Bob receives a message with Alice's digital signature on it, then:
  - Alice, and not an impersonator, sent the message (like a MAC)
  - the message has not been altered since it was sent (like a MAC), and
  - Bob can prove these facts to a third party (additional property not satisfied by a MAC).
- How do we arrange this?
  - Use similar techniques to public-key cryptography

Digital signatures are used when we don't want repudiation. It works using a public key system where the key is separated into two parts, a public and a private key.

# Making digital signatures

- Remember public-key crypto:
  - Separate keys for encryption and decryption
  - Give everyone a copy of the encryption key
  - The decryption key is private
- To make a digital signature:
  - Alice signs the message with her private **signature key**
- To verify Alice's signature:
  - Bob verifies the message with his copy of Alice's public **verification key**
  - If it verifies correctly, the signature is valid

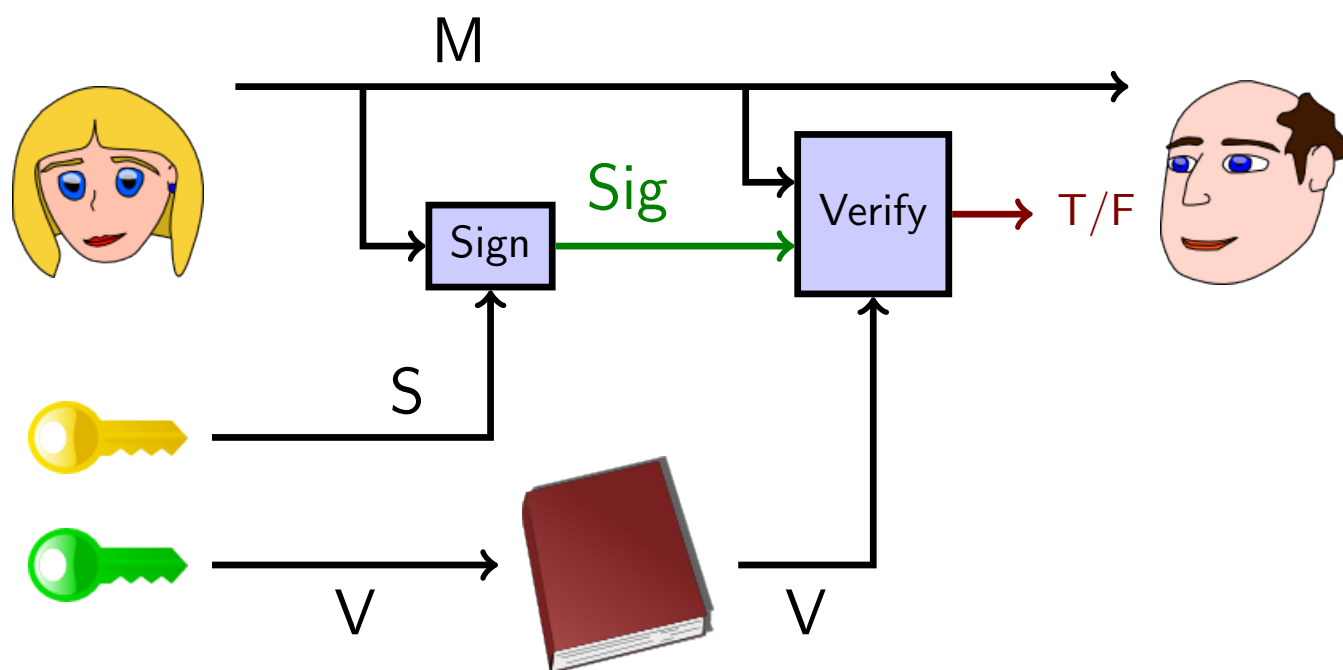


private key = signing

public key = verification

The private key is not for decryption, it is for signing things. So Alice privately signs the document and Bob can use the public key to verify the signature.

# Making digital signatures



Alice puts her secret key into some signing algorithm that also takes the message in. It outputs the signature that goes to Bob. Bob then uses the public validation key into a validation algorithm along with the signature which returns true or false if Alice is valid.

$$ab = 1 \%(p-1)(q-1)$$

$$S = M^b \% n$$

$$V = S^a == M \% n$$

S = signature, V = is valid, M = message, b = secret key, a = public key, n = private key \* public key, p and q are large prime numbers

Its very tempting to think that signing is decryption and verification is encryption because they work similarly, but they are different. There is a tone of nuance that happens in actual signing schemes to make them secure which is what makes it different from en/decryption. The textbook gets this wrong, ignore the textbook.

DSA and ECDSA are the most commonly used signature schemes. DSA is messy to talk about so we are going to look at its precursor El gammal (this led to Schnorr which led to DSA).

p = some large prime

$\alpha$  = generator for random int

$$\beta = \alpha^a \% p$$

the combination of p,  $\alpha$ , and  $\beta$  are the public key and a is the private key.

$$Sign_k(M) = (\gamma, \delta)$$

$$\gamma = \alpha^r \% p$$

$$\delta = (M - a\gamma)r^{-1} \% (p-1)$$

$$(\gamma, \delta) \rightarrow Bob$$

$$Ver_k(M, \gamma, \delta) = \beta^\gamma \gamma^\delta == \alpha^M \% p$$

# Hybrid signatures

- Just like encryption in public-key crypto, signing large messages is slow
- We can also hybridize signatures to make them faster:
  - Alice sends the (unsigned) message, and also a signature on a **hash** of the message
  - The hash is much smaller than the message, and so it is faster to sign and verify
- Remember that authenticity and confidentiality are separate; if you want both, you need to do both

Math is hard and very costly, so optimization is important.

Alice generates a hash from some hash function on her message. She then creates a signature for that hash and sends the signature and the initial message to Bob. He can then hash the message and use that to verify Alice's signature. We hash things so that the data we are running our signature and verification algorithms on is much shorter to make them much faster. We do not send the hash through to Bob because Mallory could intercept it and change it along the way.

Collisions still suck. A signature will still work for anything that has a hash collision. This means that Alice might be verifying things that she did not mean to.

# Combining public-key encryption and digital signatures

- Alice has two **different** key pairs: an (encryption, decryption) key pair and a (signature, verification) key pair
  - So does Bob
- Alice uses Bob's encryption key to encrypt a message destined for Bob
- She uses her signature key to sign the ciphertext
- Bob uses Alice's verification key to check the signature
- He uses his decryption key to decrypt the ciphertext
- Similarly for reverse direction

The recommended method is to sign then encrypt.

## Relationship between key pairs

- Alice's (signature, verification) key pair is long-lived, whereas her (encryption, decryption) key pair is short-lived
  - Gives perfect forward secrecy (see later)
- When creating a new (encryption, decryption) key pair, Alice uses her signing key to sign her new encryption key and Bob uses Alice's verification key to verify the signature on this new key
- If Alice's communication with Bob is interactive, she can use secret-key encryption and does not need an (encryption, decryption) key pair at all (see TLS or SSH)



You should periodically rotate your key pair which will give you perfect forward secrecy (if the fbi takes your hard drive they still can't decrypt everything). Everytime we rotate you publish you encryption key signed so that everyone can get your encryption key and know that it is from you.

# The Key Management Problem

- One of the hardest problems of public-key cryptography is that of **key management**
- How can Bob find Alice's verification key?
  - He can know it personally (**manual keying**)
    - SSH does this
  - He can trust a friend to tell him (**web of trust**)
    - PGP does this
  - He can trust some third party to tell him (**CA's**)
    - TLS / SSL do this

We don't care if we leak our public key, but when we get a public key we want to make sure that we are getting the correct public key and not from some man in the middle.

## Certificate authorities

- A CA is a trusted third party who keeps a directory of people's (and organizations') verification keys
- Alice generates a (signature, verification) key pair, and sends the verification key, as well as a bunch of personal information, both signed with Alice's signature key, to the CA
- The CA ensures that the personal information and Alice's signature are correct
- The CA generates a **certificate** consisting of Alice's personal information, as well as her verification key. The entire certificate is signed with the CA's signature key

# Certificate authorities

- Everyone is assumed to have a copy of the CA's verification key, so they can verify the signature on the certificate
- There can be multiple levels of certificate authorities; level  $n$  CA issues certificates for level  $n+1$  CAs
  - Public-key infrastructure (PKI)
- Need to have only verification key of root CA to verify certificate chain

When you generate a key pair you send the public key, signed, to the CA (along with identification information). The CA is supposed to check you id (like meeting in person), but usually they post a file to a domain and check that you get it. There is some higher authority that authorizes CAs (by holding their keys). Its basically a giant tree of trust. There are hundreds of root CAs

# Putting it all together

- We have all these blocks; now what?
- Put them together into **protocols**
- This is HARD. Just because your pieces all work, doesn't mean what you build out of them will; you have to *use* the pieces correctly
- Common mistakes include:
  - Using the same stream cipher key for two messages
  - Assuming encryption also provides integrity
  - Falling for replay attacks or reaction attacks
  - *LOTS* more!

# Security controls using cryptography

- Q: In what situations might it be appropriate to use cryptography as a security control?
- A: Remember that there needs to be some separation, since any secrets (like the key) need to be available to the legitimate users but not the adversaries
- In some situations, this may make secret-key crypto problematic
- If your web browser can decrypt its file containing your saved passwords, then an adversary who can read your web browser probably can, too
- Q: How is this solved in practice?



# Program and OS security

- Using secret-key crypto can be problematic for the above reason
  - But public-key is OK, if the local machine only needs access to the public part of the key
  - So only encryption and signature verification; no decryption or signing
  - E.g., apps can be installed only if digitally signed by the vendor (BlackBerry) or upgraded only if signed by the original developer (Android)
  - OS may allow execution of programs only if signed (iOS)

# Encrypted code

- There is research into processors that will only execute encrypted code
- The processor will decrypt instructions before executing them
- The decryption key is processor-dependent
- Malware won't be able to spread without knowing a processor's key
- Downsides?

# Encrypted data

- Harddrive encryption protects data when laptop gets lost/stolen
- It often does not protect data against other users who legitimately use laptop
- Or somebody installing malware on laptop
- Or somebody (maybe physically) extracting the decryption key from the laptop's memory

Some processors have security guard extensions which encrypts your instructions in such a way that only that processor can decrypt. Everytime you download code you have to tell it your processor id so that you can encrypt it. This way if malware gets on your machine it doesn't get run because it isn't encrypted. This is mostly a research idea. Most of these processors do not run everything in encrypted mode.

There is an attack known as the **cold boot attack** which is physically taking the ram from another machine and reading the values on your machine since the keys must be loaded into ram. Use liquid nitrogen on the pins of the machine to cool it down so that the values are retained in memory longer.

# OS authentication

- Authentication mechanisms often use cryptography
  - E.g., salted hashes (see Module 3)
- Unfortunately, people are bad at doing cryptography in their heads, so some mechanisms require hardware token



Photo from <http://itc.ua/>

The secure remote password protocol (challenge response) includes encryption to gain authentication. Two factor authentication is very helpful. An authentication token is really helpful here in that they just generate a key based on what time it is with some algorithm. In 2013 some hackers broke into the security company that makes the fobs and got their copy of their key value.

# Network security and privacy

- The primary use for cryptography
  - “Separating the security of the medium from the security of the message”
- Entities you can only communicate with over a network are inherently less trustworthy
  - They may not be who they claim to be

Data in transit is the main use for crypto. The data channel has to also be secure.



# Network security and privacy

- Network cryptography is used at every layer of the network stack for both security and privacy applications:
  - Link
    - WEP, WPA, WPA2
  - Network
    - VPN, IPsec
  - Transport
    - TLS / SSL, Tor
  - Application
    - ssh, Mixminion, PGP, OTR

# Link-layer security controls

- Intended to protect **local area networks**
- Widespread example: WEP (Wired Equivalent Privacy)
- WEP was intended to enforce three security goals:
  - Confidentiality
    - Prevent an adversary from learning the contents of your wireless traffic
  - Access Control
    - Prevent an adversary from using your wireless infrastructure
  - Data Integrity
- Unfortunately, **none** of these is actually enforced!

WEP is super shit. Its basically plaintext. We want to use WPA2

# WEP description

- Brief description:
- The sender and receiver share a secret  $k$ 
  - The secret  $k$  is either 40 or 104 bits long
- In order to transmit a message  $M$ :
  - Compute a checksum  $c(M)$ 
    - this does not depend on  $k$
  - Pick an IV (a random number)  $v$  and generate a keystream  $RC4(v, k)$
  - XOR  $\langle M, c(M) \rangle$  with the keystream to get the ciphertext
  - Transmit  $v$  and the ciphertext over the wireless link

A key that is 64 or 128 bits long and it uses the RC4 stream cipher where we concatenate  $v$  and  $k$  to form the keystream  $K$ . You get the cipher text through the message concatenated with the cipher of the message  $\otimes$ ed with the key stream.

One problem is that for integrity they are using a checksum where people can compute them easily.

# Problem number 0: Widely shared “secrets”



[http://www.theregister.co.uk/2014/06/25/brace\\_yourselves\\_brazil\\_dill\\_in\\_world\\_cup\\_wifi\\_spill/](http://www.theregister.co.uk/2014/06/25/brace_yourselves_brazil_dill_in_world_cup_wifi_spill/)

In this example you can see the wifi password in the background of this news broadcast.

# WEP description

- Upon receipt of  $v$  and the ciphertext:
  - Use the received  $v$  and the shared  $k$  to generate the keystream  $RC4(v, k)$
  - XOR the ciphertext with  $RC4(v, k)$  to get  $\langle M', c' \rangle$
  - Check to see if  $c' = c(M')$
  - If it is, accept  $M'$  as the message transmitted
- Problem number 1:  $v$  is 24 bits long
  - Why is this a problem?



Bob gets the keystream

$$\begin{aligned}K' &= RC4(v \parallel k) \\M' \parallel C' &= C \otimes K' \\accept = c' &== c(M')\end{aligned}$$

Assume that  $v$  is a counter (so 1, 2, 3, ...) with some max probably equal to max int which loops back to 1 when it overflows. But  $v$  is random so we can apply the birthday paradox so the number of samples required to cause a collusion is going to be the square root of int max.

## WEP data integrity

- Problem 2: the checksum used in WEP is CRC-32
  - Quite a poor choice; there's already a CRC in the protocol to detect random errors, and a CRC can't help you protect against malicious errors.
- The CRC has two important properties:
  - It is independent of  $k$  and  $v$
  - It is **linear**:  $c(M \text{ XOR } D) = c(M) \text{ XOR } c(D)$
- Why is linearity a pessimal property for your integrity mechanism to have when used in conjunction with a stream cipher?

WEP uses the cyclic redundancy check to calculate the checksum which kinda sucks. Checksums are not meant to protect against attackers.

An adversary picks some  $M' = M \otimes \delta$  and gets the cipher text  $c(M') = c(M' \otimes \delta) = c(M') \otimes c(\delta)$ . The attacker wants some  $C' = (M' \parallel c(M')) \otimes K$  but they do not know what K is.

Do some math:

$$\begin{aligned} C' &= (M' \parallel c(M')) \otimes K \\ C' &= (M \otimes \delta \parallel c(M) \otimes x(\delta)) \otimes KC' = K \otimes (M \otimes \delta \parallel c(M) \otimes x(\delta))C' = K \otimes ((M \parallel c(M)) \otimes (\delta \parallel c(\delta))) \end{aligned}$$

We see that  $K \otimes ((M \parallel c(M)))$  is just the key stream so an attacker can just take the first chunk of a message sent and build her message out of that in such a way that the integrity check still passes.

## WEP access control

- What if the adversary wants to inject a new message  $F$  onto a WEP-protected network?
- All he needs is a single plaintext/ciphertext pair
- This gives him a value of  $v$  and the corresponding keystream  $RC4(v, k)$
- Then  $C' = \langle F, c(F) \rangle \text{ XOR } RC4(v, k)$ , and he transmits  $v, C'$
- $C'$  is in fact a correct encryption of  $F$ , so the message must be accepted

Say the attacker gets the cipher text and the plain text pair, they can use it to calculate the key string.

$$RC4 = C \otimes (M \parallel c(M)) = K$$

Not to send some made up message F we can calculate the checksum because we now have some K to use. WEP allows us to reuse key streams since collisions are common. They do not check for reuse. So once we get a keystream for any message we can use it to encrypt all messages.

# WEP authentication protocol

- How did the adversary get that single plaintext/ciphertext pair required for the attack on the previous slide?
  - Problem 3: It turns out the authentication protocol gives it to the adversary **for free**!
- This is a major disaster in the design!
- The authentication protocol (described on the next slide) is supposed to prove that a certain client knows the shared secret  $k$
- But if I watch you prove it, I can turn around and execute the protocol myself!
  - “What’s the password?”

When you first connect to an access point it sends you a plaintext value  $M$  which is a random challenge. You encrypt it and send it back along with some initialization vector  $v$ . The attacker just has to get  $v$  and they can send as many encrypted messages as they want.

# WEP authentication protocol

- Here's the authentication protocol:
  - The access point sends a challenge string to the client
  - The client sends back the challenge, WEP-encrypted with the shared secret  $k$
  - The wireless access point checks if the challenge is correctly encrypted, and if so, accepts the client
- So the adversary has just seen both the plaintext and the ciphertext of the challenge
- Problem number 4: this is enough not only to inject packets (as in the previous attack), but also to execute the authentication protocol himself!



# WEP decryption

- Somewhat surprisingly, the ability to modify and inject packets also leads to ways the adversary can **decrypt** packets!
  - The access point knows  $k$ ; it turns out the adversary can trick it into decrypting the packet for him and telling him the result.
- Note that none of the attacks so far:
  - Used the fact that the stream cipher was RC4 specifically
  - Recovered  $k$

# Recovering a WEP key

- Since 2002, there have been a series of analyses of RC4 in particular
  - Problem number 5: it turns out that when RC4 is used with similar keys, the output keystream has a subtle weakness
    - And this is (often) how WEP uses RC4!
- These observations have led to programs that can recover either a 104-bit or 40-bit WEP key in **under 60 seconds**, most of the time
  - See the optional reading for more information on this

# Replacing WEP

- Wi-fi Protected Access (WPA) was rolled out as a short-term patch to WEP while formal standards for a replacement protocol (IEEE 802.11i, later called WPA2) were being developed
- WPA:
  - Replaces CRC-32 with a real MAC (here called a MIC to avoid confusion with a Media Access Control address)
  - IV is 48 bits
  - Key is changed frequently (TKIP)
  - Ability to use 802.1x authentication server
    - But maintains less-secure PSK (Pre-Shared Key) mode for home users
  - Able to run on most older WEP hardware

# Replacing WEP

- The 802.11i standard was finalized in 2004, and the result (called WPA2) has been required for products calling themselves “Wi-fi” since 2006
- WPA2:
  - Replaces the RC4 and MIC algorithms in WPA with the CCM authenticated encryption mode (using AES)
  - Considered strong, except in PSK mode
    - Dictionary attacks still possible

Dictionary attacks are always possible for access control. So if you run one against WPA2 you can get the key and decrypt data. There are some protocols that will protect against dictionary attacks (ex Diffie-Hellman).

# Network-layer security

- Suppose every link in our network had strong link-layer security
- Why would this not be enough?
- We need security **across** networks
  - Ideally, **end-to-end**
- At the network layer, this is usually accomplished with a Virtual Private Network (VPN)

We assume that any physical network is secure, but there are many networks that need to talk to each other (think of a company with many locations). This is where network layer security protocols come in.

# Virtual Private Networks

- Connect two (or more) networks that are physically isolated, and make them appear to be a single network
  - Alternately: connect a single remote host (often a laptop) to one network
- Goal: adversary between the networks should not be able to read or modify the traffic flowing across the VPN
  - But DoS and some traffic analysis still usually possible



We use VPNs to connect multiple networks that are in different physical locations so that they seem to be a single network. Alice wants to connect so she just connects to a VPN which directs the connection to a physical location but she just sees it as a single network. An attacker who is sitting in the cloud can see the traffic but you don't want them to know what you are doing so you go through a VPN which can hide stuff.

# Setting up a VPN

- One host on each side is the **VPN gateway**
  - Could be the firewall itself, or could be in DMZ
  - In the laptop scenario, it will of course be the laptop itself on its side
- Traffic destined for the “other side” is sent to the local VPN gateway
- The local VPN gateway uses cryptography (encryption and integrity techniques) to send the traffic to the remote VPN gateway
  - Often by **tunnelling** (see next slide)
- The remote VPN gateway decrypts the messages and sends them on to their appropriate destinations

VPN gateways are hosts on each network that do the special calculations required for routing traffic to a different network through the VPN (address translation and such).

# Tunnelling

- Tunnelling is the sending of messages of one protocol inside (that is, as the payload of) messages of another protocol, out of their usual protocol nesting sequence
- So TCP-over-IP **is not** tunnelling, since you're supposed to send TCP (a transport protocol) over IP (a network protocol; one layer down in the stack)
- But IP-over-TCP **is** tunnelling (going up the stack instead of down), as are IP-over-IP (same place in the stack), and PPP (a link layer protocol; bottom of the stack) over DNS (an application layer protocol; top of the stack)

To get free wifi on a network that uses a captive portal. Usually the way they implement this is by catching your initial http request (instead of the dns request) and directing it to their site. You can use tunneling to take all the packets you want to send and wrap them into the application layer to turn it into a dns packet. You can send these to a home server which unpacks them, sends it, gets the response, turns it into a dns packet, and responds with it to you. Now you can use their wifi.

In tunnel mode everything is encrypted and authenticated. It then adds a new IP header where the source is the VPN gateway of the sender and the destination is the VPN gateway of the destination. These gateways do this packaging and unpackaging when packets go through them.

The data is encrypted as it goes so listeners cannot tell what data is being sent, only that two people are talking to each other

# IPsec

- One standard way to set up a VPN is by using IPsec
- Many corporate VPNs use this (open) protocol
- Two modes:
  - **Transport** mode
    - Useful for connecting a single laptop to a home network
    - Only the contents of the original IP packet are encrypted and authenticated
  - **Tunnel** mode
    - Useful for connecting two networks
    - The contents **and the header** of the original IP packet are encrypted and authenticated; result is placed inside a new IP packet destined for the remote VPN gateway

The previous slide basically described what happens in tunnel mode. Transport mode works for remote machines connecting over the internet to the VPN. Alice, in this case, has to run a special client to do so. The TCP layer and transport layers are encrypted, but the header is not, so the source and destination are still in plain text. The packet gets to the VPN gateway which decrypts the payload and forwards it. So someone listening in the cloud knows who is talking to who in the network.

## Other styles of VPNs

- In addition to IPsec, there are a number of other standard ways to set up a VPN
- Microsoft's PPTP was an older protocol
  - It had about as many design flaws as WEP
  - Most users now migrating to IPsec
- VPNs based on ssh
  - Tunnel PPP over ssh
    - That is, IP-over-PPP-over-ssh-over-TCP-over-IP
    - Some efficiency concern, but extremely easy to set up on a standard Unix/Linux box
  - OpenSSH v4 supports IP-over-SSH tunnelling directly



PPTP has a tone of design flaws like WEP so most people do not use it. The most common one for personal use is Open VPN. You can tunnel over ssh

# Transport-layer security and privacy

- Network-layer security mechanisms arrange to send individual IP packets securely from one network to another
- Transport-layer security mechanisms transform arbitrary TCP connections to add security
  - And similarly for “privacy” instead of “security”
- The main transport-layer security mechanism:
  - TLS (formerly known as SSL)
- The main transport-layer privacy mechanism:
  - Tor

Transport layer security is used for tons of stuff. TLS provides everything except privacy and TOR is meant for privacy.

# TLS / SSL

- In the mid-1990s, Netscape invented a protocol called Secure Sockets Layer (SSL) meant for protecting HTTP (web) connections
  - The protocol, however, was general, and could be used to protect **any** TCP-based connection
  - HTTP + SSL = **HTTPS**
- Historical note: there was a competing protocol called S-HTTP. But Netscape and Microsoft both chose HTTPS, so that's the protocol everyone else followed
- SSL went through a few revisions, and was eventually standardized into the protocol known as **TLS** (Transport Layer Security, imaginatively enough)

Netscape set a bunch of standards by inventing stuff they wanted to have and being so common that future browsers built on that.

# TLS at a high level

- Client connects to server, indicates it wants to speak TLS, and which **ciphersuites** it knows
- Server sends its certificate to client, which contains:
  - Its host name
  - Its verification key
  - Some other administrative information
  - A signature from a Certificate Authority (CA)
- Server also chooses which ciphersuite to use

Alice sends a bunch of cipher suites (a list of everything the browser supports). The server responds with its certificate and choice of cipher suite. The certificate comes from a certificate authority who's decryption key is installed in her browser. Alice then verifies the certificate.

Occasionally people install the wrong certificate on the wrong server. Its actually kinda hard to do on purpose. Most likely the website was hosted by cloud flair who makes certificates for websites, so they probably made the mistake.

## TLS at a high level (cont.)

- Client validates server's certificate
  - Is its signature from a CA whose public key is embedded in the client (i.e., browser)?
  - Does the host name in the certificate match the host name of the web site that client wants to download?
- Client and server run a key agreement protocol to establish keys for symmetric encryption and MAC algorithms from the chosen ciphersuite
  - Server signs its protocol messages with its signature key
- Communication now proceeds using chosen symmetric encryption and MAC algorithms



Alice generates a master secret (some large value to be used as a key) then generates a hash of the master key and a salt which is the MAC key and then generates another one with a different salt to be the AES key. She then uses the server's public encryption key that she uses to encrypt the master key and sends it to the server. The server then decrypts this and hashes it with salts to get the MAC and AES keys.

This is a super simple version of what is actually used (called a key agreement protocol). Diffie-Hellman is the most commonly used one, but we'll cover it later.

# Certificate validation

This server could not prove that it is **convention.gop**; its security certificate is from **funyuns.com**. This may be caused by a misconfiguration or an attacker intercepting your connection.

[Proceed to convention.gop \(unsafe\)](#)

The messages being send by TLS are in plaintext so Mallory can sit in the middle and can alter the list of cipher suites used by Alice to be the weakest one. This allows Mallory to attack this. This is called a **downgrade attack**.

One problem with client authentication is that you have to issue certificates to every client which is annoying. Also key management is hard and certificates are expensive.

# Security properties provided by TLS

- Server authentication
- Message integrity
- Message confidentiality
- Client authentication (optional)

We did a really good job getting users to want ssl (look for the lock). It used to be common to get an SSL error when connecting to a website. This caused warning fatigue on users to just click ignore the warning.

For a long time there was a website for administrators at UW to use. The certificate they were using was self signed so every time they connected they would get the SSL warning. So part of the actual training for admins was to click ignore on that.

Now-a-days they made the warning much harder to get around so people tend to just ignore the website. This gave websites a strong incentive to get their shit secure. There is even a project that will give you a TLS certificate for free.

# The success of TLS

- Though designed as a security mechanism, TLS (including SSL) has become the most successful **Privacy Enhancing Technology (PET)** ever
- Why?
  - It comes with your browser
    - Which encouraged web server operators to bother paying \$\$ for their certificates
  - It just works, without you having to configure anything
  - Most of the time, it even protects the privacy of your communications
    - Increasingly important due to the success of WiFi

So far encryption doesn't protect the metadata of our data. It doesn't protect who sent it to who. This is because we need to be able to route messages along the way. This is where TOR comes into play.

# Privacy Enhancing Technologies

- So far, we've only used encryption to protect the **contents** of messages
- But there are other things we might want to protect as well!
- We may want to protect the **metadata**
  - Who is sending the message to whom?
  - If you're seen sending encrypted messages to Human Rights Watch, bad things may happen
- We may want to hide the **existence** of the message
  - If you're seen sending encrypted messages at all, bad things may happen



TOR works on the transport layer. It has about 2 million users. You can download the TOR browser bundle (its a version of firefox that uses TOR). The most common number of connections is HTTP. The most traffic over TOR is from bit torrent.

# Tor

- **Tor** is another successful privacy enhancing technology that works at the transport layer
  - Hundreds of thousands of users
  - March 2015:  $\approx$  2 million users
- Normally, a TCP connection you make on the Internet automatically reveals your IP address
  - Why?
- Tor allows you to make TCP connections **without** revealing your IP address
- It's most commonly used for HTTP (web) connections

# How Tor works


- Scattered around the Internet are about 7,000 Tor **nodes**, also called **Onion Routers**
- Alice wants to connect to a web server without revealing her IP address



The Tor Browser uses onion encryption (lots of layers). There are nodes scattered around the world.


# How Tor works

- Alice picks one of the Tor nodes (n1) and uses public-key cryptography to establish an encrypted communication channel to it (much like TLS)




# How Tor works

- Alice tells  $n_1$  to contact a second node ( $n_2$ ), and establishes a new encrypted communication channel to  $n_2$ , tunnelled within the previous one to  $n_1$



# How Tor works

- Alice tells  $n_2$  to contact a third node ( $n_3$ ), and establishes a new encrypted communication channel to  $n_3$ , tunnelled within the previous one to  $n_2$




The first node that Alice hits is the entry guard. She shares a key with it  $k_1$  via the normal TLS key agreement method. She then tells  $n_1$  to contact a second node,  $n_2$  on my behalf.  $n_2$  is a middle node. Alice sends key  $k_2$  to  $n_2$  and so on. Alice gets to pick all of the nodes that she wants to pass shit along. None of the intermediate nodes can read the messages that she sends to the others. Evnetually (in her example 3 nodes later) she reaches an exit node. After this the packet exits the tor network. Only the nodes along this circuit have her keys.



# How Tor works

- And so on, for as many steps as she likes (usually 3)
- Alice tells the last node (within the layers of tunnels) to connect to the website



Say Alice wants to send message  $M$  to webserver. She takes this message and encrypts it in reverse order (so encrypt with  $k_3$  then  $k_2$ , and so on). Each node decrypts a single layer as it goes.

## Sending messages with Tor

- Alice now shares three secret keys:
  - $K1$  with  $n1$
  - $K2$  with  $n2$
  - $K3$  with  $n3$
- When Alice wants to send a message  $M$ , she actually sends  $E_{K1}(E_{K2}(E_{K3}(M)))$
- Node  $n1$  uses  $K1$  to decrypt the outer layer, and passes the result  $E_{K2}(E_{K3}(M))$  to  $n2$
- Node  $n2$  uses  $K2$  to decrypt the next layer, and passes the result  $E_{K3}(M)$  to  $n3$
- Node  $n3$  uses  $K3$  to decrypt the final layer, and sends  $M$  to the website

Basically it just does the exact same thing in reverse.

## Replies in Tor

- When the website replies with message  $R$ , it will send it to node  $n3$ 
  - Why?
- Node  $n3$  will **encrypt**  $R$  with  $K3$  and send  $E_{K3}(R)$  to  $n2$
- Node  $n2$  will encrypt that with  $K2$  and send  $E_{K2}(E_{K3}(R))$  to  $n1$
- Node  $n1$  will encrypt that with  $K1$  and send  $E_{K1}(E_{K2}(E_{K3}(R)))$  to Alice
- Alice will use  $K1$ ,  $K2$ , and  $K3$  to decrypt the layers of the reply and recover  $R$

It is possible to set up an exit node that can watch all the plaintext exiting the node. Make sure you use https or some other secure form of connection.

## Who knows what?

- Notice that node n1 knows that Alice is using Tor, and that her next node is n2, but does not know which website Alice is visiting
- Node n3 knows some Tor user (with previous node n2) is using a particular website, but doesn't know who
- The website itself only knows that it got a connection from Tor node n3
- **Note:** the connection between n3 and the website is **not encrypted**! If you want encryption as well as the benefits of Tor, you should use end-to-end encryption **in addition**
  - Like HTTPS

# Anonymity vs. pseudonymity

- Tor provides for **anonymity** in TCP connections over the Internet, both **unlinkably** (long-term) and **linkably** (short-term)
- What does this mean?
  - There's no long-term identifier for a Tor user
  - If a web server gets a connection from Tor today, and another one tomorrow, it won't be able to tell whether those are from the same person
  - But two connections in quick succession from the same Tor node are more likely to in fact be from the same person



# The Nymity Slider

- We can place transactions (both online and offline) on a continuum according to the level of **nymity** they represent:
  - **Verinymity**
    - Government ID, SIN, credit card #, address
  - **Persistent pseudonymity**
    - Noms de plume, many blogs
  - **Linkable anonymity**
    - Prepaid phone cards, loyalty cards
  - **Unlinkable anonymity**
    - Cash payments, Tor

# The Nymity Slider

- If you build a system at a certain level of nymity, it's **easy** to modify it to have a higher level of nymity, but **hard** to modify it to have a lower level
- For example:
  - It's easy to add a loyalty card to a cash payment, or a credit card to a loyalty card
  - It's hard to remove identity information if you're paying by credit card
- The lesson: design systems with a low level of nymity fundamentally; adding more is easy

# Application-layer security and privacy

- TLS can provide for encryption at the TCP socket level
  - “End-to-end” in the sense of a network connection
  - Is this good enough? Consider SMTPS (SMTP/email over TLS)
- Many applications would like true end-to-end security
  - Human-to-human would be best, but those last 50 cm are really hard!
  - We usually content ourselves with desktop-to-desktop
- We’ll look at three particular applications:
  - Remote login, email, instant messaging

## Secure remote login (ssh)

- You're already familiar with this tool for securely logging in to a remote computer (the ugster machines)
- Usual usage (simplified):
  - Client connects to server
  - Server sends its verification key
    - The client **should** verify that this is the correct key
  - Client and server run a key agreement protocol to establish session keys, server signs its messages
    - All communication from here on in is encrypted and MACd with the session keys
  - Client authenticates to server
  - Server accepts authentication, login proceeds (under encryption and MAC)

# Authentication with ssh

- There are two main ways to authenticate with ssh:
  - Send a password over the encrypted channel
    - The server needs to know (a hash of) your password
  - Sign a random challenge with your private signature key
    - The server needs to know your public verification key
- Which is better? Why?

# Pretty Good Privacy

- The first popular implementation of public-key cryptography.
- Originally made by Phil Zimmermann in 1991
  - He got in a lot of trouble for it, since cryptography was highly controlled at the time.
  - But that's a whole 'nother story. :-)
- Today, there are many (more-or-less) compatible programs
  - GNU Privacy Guard (gpg), Hushmail, etc.

# Pretty Good Privacy

- What does it do?
  - Its primary use is to protect the contents of email messages
- How does it work?
  - Uses public-key cryptography to provide:
    - Encryption of email messages (using hybrid encryption)
    - Digital signatures on email messages (hash-then-sign)

# Recall

- In order to use public-key encryption and digital signatures, Alice and Bob must each have:
  - A public encryption key
  - A private decryption key
  - A private signature key
  - A public verification key



# Sending a message

- To send a message to Bob, Alice will:
  - Write a message
  - Sign it with her own signature key
  - Encrypt both the message and the signature with Bob's public encryption key
- Bob receives this, and:
  - Decrypts it using his private decryption key to yield the message and the signature
  - Uses Alice's verification key to check the signature

# Back to PGP

- PGP's main functions:
  - Create these four kinds of keys
    - encryption, decryption, signature, verification
  - Encrypt messages using someone else's encryption key
  - Decrypt messages using your own decryption key
  - Sign messages using your own signature key
  - Verify signatures using someone else's verification key
  - Sign other people's keys using your own signature key (see later)

## HOW TO USE PGP TO VERIFY THAT AN EMAIL IS AUTHENTIC:

LOOK FOR THIS  
TEXT AT THE TOP.



IF IT'S THERE, THE EMAIL IS PROBABLY FINE.

to be extra safe, check that there's a big block of jumbled characters at the

# Obtaining keys

- Earlier, we said that Alice needs to get an authentic copy of Bob's public key in order to send him an encrypted message
- How does she do this?
  - In a secure way?
- Bob could put a copy of his public key on his webpage, but this isn't good enough to be really secure!
  - Why?

# Verifying public keys

- If Alice knows Bob personally, she could:
  - Download the key from Bob's web page
  - Phone up Bob, and verify she's got the right key
  - Problem: keys are big and unwieldy!

```
mQGIBDi5qEURBADitpDzvzW+9lj/zYgK78G3D76hvvtIT6gpTIlwg6WIJNLKJat
01yNpMIYNvpwi7EUd/lSN16t1/A022p7s7bDbE4T5NJda0IOAgWe0Z/plIJC4+o2
tD2RNuSkwDQcxzm8KUNZ0Jla4Lvgrkm/oUubxyeY5omus7hcfNrB0wjC1wCg4Jnt
m7s3eNfMu72Cv+6FzBgFog8EANirkNdC1Q8oSMDihWj1ogiWbBz4s6HMxzAaqNf/
rCJ9qoK5SLFeoB/r5ksRWty9QKV4VdhhCIy1U2B9tSTlEPYXJHQPZ3mwCxUnJpGD
8UgFM5uKXaEq2pwpArTm367k0tTpMQgXAN2HwiZv//ahQXH4ov30kBBVL5VFxMUL
UJ+yA/4r5HLTpP2SbbqtPWdeW7uDwhe2dTqffAGuf0kuCpHwCTAhr83ivXzT/7OM
```

# Fingerprints

- Luckily, there's a better way!
- A **fingerprint** is a cryptographic hash of a key
- This, of course, is *much shorter*:
  - B117 2656 DFF9 83C3 042B C699 EB5A 896A 2898 8BF5
- Remember: there's no (known) way to make two different keys that have the same fingerprint, provided that we use a collision-resistant hash function

# Fingerprints

- So now we can try this:
  - Alice downloads Bob's key from his webpage
  - Alice's software calculates the fingerprint
  - Alice phones up Bob, and asks him to read his key's actual fingerprint to her
  - If they match, Alice knows she's got an authentic copy of Bob's key
- That's great for Alice, but what about Carol, who doesn't know Bob
  - At least not well enough to phone him

## Signing keys

- Once Alice has verified Bob's key, she uses her signature key to sign Bob's key
- This is effectively the same as Alice signing a message that says "I have verified that the key with fingerprint B117 2656 DFF9 83C3 042B C699 EB5A 896A 2898 8BF5 really belongs to Bob"
- Bob can attach Alice's signature to the key on his webpage



# Key signing parties



Alice generates a signature on Bob's public key to say that she totally verified Bob's fingerprint. Bob takes that signature and put it in a public place so that anyone who trusts Alice will now trust Bob,

# Web of Trust

- Now Alice can act as an introducer for Bob
- If Carol doesn't know Bob, but does know Alice (and has already verified Alice's key, and trusts her to introduce other people):
  - she downloads Bob's key from his website
  - she sees Alice's signature on it
  - she is able to use Bob's key without having to check with Bob personally
- This is called the Web of Trust, and the PGP software handles it mostly automatically

There are three kinds of keys, Trust, Partially Trusted, and Untrusted. Alice assigns one of these to each key she adds to her key ring.

# So, great!

- So if Alice and Bob want to have a private conversation by email:
  - They each create their sets of keys
  - They exchange public encryption keys and verification keys
  - They send signed and encrypted messages back and forth
- Pretty Good, no?

## Plot twist

- Suppose (encrypted) communications between Alice and Bob are recorded by the “bad guys”
  - criminals, competitors, subpoenaed by the RCMP
- Later, Bob’s computer is stolen by the same bad guys
- Or just broken into
  - Virus, trojan, spyware
- All of Bob’s key material is discovered
  - Oh, no!

## The bad guys can...

- Decrypt past messages
- Learn their content
- Learn that Alice sent them
- And have a mathematical **proof** they can show to anyone else!
- How private is that?

# What went wrong?

- Bob's computer got stolen?
- How many of you have never...
  - Left your laptop unattended?
  - Not installed the latest patches?
  - Run software with a remotely exploitable bug?
- What about your friends?



# What really went wrong

- PGP creates lots of incriminating records:
  - Key material that decrypts data sent over the public Internet
  - Signatures with proofs of who said what
- Alice had better watch what she says!
  - Her privacy depends on Bob's actions

# Casual conversations

- Alice and Bob talk in a room
- No one else can hear
  - Unless being recorded
- No one else knows what they say
  - Unless Alice or Bob tells them
- No one can prove what was said
  - Not even Alice or Bob
- These conversations are “off-the-record”

# We like off-the-record conversations

- Legal support for having them
  - Illegal to record other people's conversations without notification
- We can have them over the phone
  - Illegal to tap phone lines
- But what about over the Internet?

# Crypto tools

- We have the tools to do this
  - We've just been using the wrong ones
  - (when we've been using crypto at all)
- We want perfect forward secrecy
- We want deniable authentication

## Perfect forward secrecy

- Future key compromises should not reveal past communication
- Use secret-key encryption with a short-lived key (a **session key**)
- The session key is created by a modified Diffie-Hellman protocol
- Discard the session key after use
  - Securely erase it from memory
- Use long-term keys only to authenticate the Diffie-Hellman protocol messages
- Q: Why does this approach not have the very same forward secrecy problem as PGP?

We want to try to share a key with perfect forward secrecy. The Diffie-Hellman protocol is the best way to do this.

Get a prime number  $p$  and a generator  $g$  of prime integers. Alice picks a random  $a$  that remains private and sends  $g^a \bmod p$  to Bob. Bob picks a random  $b$  that remains private and sends  $g^b \bmod p$  to Alice. Alice then generates a key  $K = (g^b)^a \bmod p$  and Bob generates  $K = (g^a)^b \bmod p$  resulting in a shared key. Currently there is no known way for Eve to calculate  $a$  or  $b$ . So far incalculable.

# Deniable authentication

- Do **not** want digital signatures
  - Non-repudiation is great for signing contracts, but undesirable for private conversations
- But we **do** want authentication
  - We can't maintain privacy if attackers can impersonate our friends
- Use Message Authentication Codes
  - We talked about these earlier

To prevent man in the middle attacks we sign the bits getting passed back and forth. We still want deniability though (so that you cannot say for certain which trusts source sent it). So signatures aren't the best. So we use MACs



## No third-party proofs

- Shared-key authentication
  - Alice and Bob have the same MK
  - MK is required to compute the MAC
  - How is Bob assured that Alice sent the message?
- Bob cannot prove that Alice generated the MAC
  - He could have done it, too
  - Anyone who can verify can also forge
- This gives Alice a measure of deniability

Bob is assured that only Alice sent the message since no one else had access to the MAC. Here we cannot deny that we have initiated a conversation, but by using the MAC we get deniability on who sent which message.

## Using these techniques

- Using these techniques, we can make our online conversations more like face-to-face “off-the-record” conversations
- But there’s a wrinkle:
  - These techniques require the parties to communicate interactively
  - This makes them unsuitable for email
  - But they’re still great for instant messaging!

# Off-the-Record Messaging

- Off-the-Record Messaging (OTR) is software that allows you to have private conversations over instant messaging, providing:
- Confidentiality
  - Only Bob can read the messages Alice sends him
- Authentication
  - Bob is assured the messages came from Alice

We like OTR. It gives confidentiality, and perfect forward secrecy. It also uses key rotation. The section of time using the same key is an epoch so if a hacker somehow gets the key, they can only read what was said during that epoch.

# Off-the-Record Messaging

- Perfect Forward Secrecy
  - Shortly after Bob receives the message, it becomes unreadable to anyone, anywhere
- Deniability
  - Although Bob is assured that the message came from Alice, he can't convince Charlie of that fact
  - Also, Charlie can create forged transcripts of conversations that are every bit as accurate as the real thing

At the end of an epoch the current MAC key is published so that we have deniability (the attacker can forge old transcripts, but not effect the current one). Alice and Bob are already using a new key so the attacker cannot pretend to be them.

# Signal Protocol

- Signal is an app for iOS and Android
  - Original protocol based on OTR and used for encrypted SMS
- The Signal Protocol is now used by WhatsApp, Google Allo, and Facebook Messenger
- Provides forward secrecy, improved deniability



Signal doesn't even know what you are talking about.

# Signal Protocol

- Perfect forward secrecy
  - Similar to OTR, uses a “ratchet” technique to constantly rotate session keys
- Deniability
  - Uses “Triple Diffie-Hellman” deniable authenticated key exchange (DAKE)
  - Anyone can forge a conversation between Alice and Bob using only their public keys

Signal uses the triple Diffie-Hellman protocol. Alice has a long term key. She generates an effemeral (short term session) key for her public and private keys

- long term private =  $A$
- short term private =  $a$
- long term public =  $g^A$
- short term public =  $g^a$

Bob also generates these keys using  $b$  and  $B$ .

So Alice sends her two public keys to Bob and he sends his back, using this they create keys by hashing a concatenation of three separate keys (hence the triple in the name of this protocol).

Now when we generate the shared key we do  $K = h(g^{ab} || g^{Ab} || g^{aB})$

Anyone can forge a conversation using the public keys by generating random private keys for each user and just making stuff up. This will look like a conversation between Alice and Bob. It cannot be used as an attack because you have to make up the private keys for both.

# Recap

- Basics of cryptography
- Secret-key encryption
- Public-key encryption
- Integrity
- Authentication
- Security controls using cryptography
- Link-layer security
- Network-layer security
- Transport-layer security and privacy
- Application-layer security and privacy